

Course Overview

Prof. Hyuk-Yoon Kwon

OS
RAM
CPU Architecture.
Bit level.

Contents

■ Part1: Course Overview

- Course Policies
- Course Theme
- Why This Course Important? Five realities!

■ Part2: The Overview of Computer Systems

Course Objective

- In this course, we focus on **what really happens when your programs run**. The aim of this course is to explain the enduring concepts **underlying all computer systems**, and to show you the concrete ways that these ideas affect the **correctness, performance, and utility** of your application programs.

Course Description

- This course deals with the basic concepts underlying all computer systems **from a programmer's perspective**. It describes how application programmers can use their knowledge of a system **to write better programs**. Specifically, this course will cover **how to represent and manipulate data in computers**, machine-level representation of programs, processor architecture, memory hierarchy, and virtual memory.

Lecture notes and textbooks provide many useful program samples **written in C programming language**, which is still the representative language to implement system-layer applications. During the entire course, it is recommended to be familiar with C language. For the students who are not familiar with C, we will cover the basics of C in the beginning of the course.

Course Policy: Lectures

■ Flipped learning

- (Online lecture contents) The contents will be uploaded **by the previous Tuesday**
- (Offline lecture) **Every Monday 11:00 AM ~** *1시간 정도 대면수업 예정!*
 - Summary (will be given in English)
 - Q&A (Korean questions are also allowed)

■ Lecture place

- Laboratory: **Frontier 511**

Course Policy: Grading

■ Exams (70%)

- Mid-term (30%)
- Final (40%)

■ Assignments (30%)

- Homework (10%) → 374
- Projects (20%) 2 Team Projects (2-3 members)
↳ Making Program.

Course Policy: Homework

■ Penalty for cheating

- The lowest grade

■ What is cheating?

- **Sharing code**: by copying, retyping, looking at, or supplying a file
- **Searching the Web for solutions**
- Helping your friend to write a lab, line by line

■ What is NOT cheating?

- Helping others with high-level design issues
- Explaining how to use systems or tools

Course Assumption

■ Required: Basic programming knowledge

- Any kind of programming languages - Java, C/C++, or Python

■ C Programming Language

- Lecture notes and textbooks provide many useful program samples written in C
- C is the representative language to implement system-layer applications
- For the students who are not familiar with C, we will cover the basics of C in the next lecture

```
1  #include <stdio.h>
2
3  int main()
4  {
5      printf("hello, world\n");
6      return 0;
7  }
```

code/intro/hello.c

Figure 1 A typical code example.

Environments for Practices

■ Recommended environment: Linux

- Sample codes in lecture notes and textbooks are written in Linux

■ Possible environments

- Install Linux OS: any kind of Linux OS is allowed
 - Most Linux systems are open sourced
 - Example: <https://www.ubuntu.com/> (you can download here)
- Install Virtual Machines on Windows
 - **Vmware (recommended)** or Virtual box for Windows
 - You can install Linux on virtual machines in Windows
- Or
 - Cloud service provided by Amazon or Google
 - Any kind of environments to use Linux systems are possible

■ The easiest recommended way will be provided by manual

Textbooks

■ Randal E. Bryant and David R. O'Hallaron,

- *Computer Systems: A Programmer's Perspective*, Third Edition (CS:APP3e), Pearson, 2016
- <http://csapp.cs.cmu.edu>
- This book really matters for the course!
 - How to solve labs
 - Practice problems typical of exam problems

■ Brian Kernighan and Dennis Ritchie,

- *The C Programming Language*, Second Edition, Prentice Hall, 1988
- Still the best book about C, from the originators

Getting Help

■ E-class: <http://eclass.seoultech.ac.kr>

- Complete schedule of lectures, exams, and assignments
- Lecture slides, assignments, exams, solutions
- Clarifications to assignments
- Q&A

■ Office hours

- Monday, 2:00-4:00pm, Frontier 614
- You can schedule 1:1 appointments in advance

Course Schedule

가. 2022학년도 2학기 공휴일 현황

수업주차	공휴일자	공휴일명	비고
1주차 ['22. 9. 5.(월) ~ 9. 9.(금)]	'22. 9. 9.(금)	추석	
2주차 ['22. 9.12.(월) ~ 9.16.(금)]	'22. 9.12 (월)	대체공휴일	추석
5주차 ['22.10. 3.(월) ~10. 7.(금)]	'22.10. 3.(월)	개천절	
6주차 ['22.10.10.(월) ~10.14.(금)]	'22.10.10.(월)	대체공휴일	한글날
12주차 ['22.11.21.(월) ~11.25.(금)]	'22.11.21.(월)~11.22.(화)	논술고사일	임시휴업일

9/12 (추석)
10/3 (개천절)
10/10 (한글날)

Week	Contents	Offline (Flipped Learning)
1	Overview	Introduction of course
2	Introduction of C Programming Languages	No offline lecture
3	Information Storage	Summary for 1 st and 2 nd lecture
4	Integer/Floating Representation	Summary for 3 rd and 4 th lecture
5	Machine-Level Programming: Basics	No offline lecture
6	Machine-Level Programming: Controls	No offline lecture
7	Machine-Level Programming: Procedures	Summary for 5 th and 6 th lecture
8	Mid-Term Exam	
9	1 st Project Presentation	

1-6 week.

11/21 (논술)

Week	Contents	Offline (Flipped Learning)
10	Machine-Level Programming: Data	
11	Machine-Level Programming: Advanced	No offline lecture
12	Memory Hierarchy	
13	Cache Memories and Virtual Memories	
14	2 nd Project Presentation	
15	Final Exam	

Contents

■ Part1: Course Overview

- Course Policies
- Course Theme
- Why This Course Important? Five realities!

■ Part2: The Overview of Computer Systems

Course Theme: Abstraction Is Good But Don't Forget Reality

■ Abstraction is important

- Asymptotic analysis
 - If $f(n) = n^2 + 3n$, then as n becomes very large, $3n$ becomes insignificant compared to n^2

■ The abstraction has limits


- Especially in the presence of bugs
- Need to understand **details of underlying implementations**

■ Useful outcomes from this course

- Become more **effective programmers**
 - Able to find and eliminate bugs efficiently
 - Able to understand for program performance
- Become more **effective IT consultant**

Great Reality #1: Ints are not Integers?

■ Example : Is $x^2 \geq 0$?



Source: xkcd.com/571

```
int a = 40000 * 40000; // 1600000000
int b = 50000 * 50000; // 2500000000 ??
```


Computer Arithmetic

■ Does not generate random values

- Arithmetic operations have important mathematical properties

■ Cannot assume all “usual” mathematical properties

- Due to **finiteness of representations**

■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

Great Reality #2: You've Got to Know Assembly

■ Chances are, you'll never write programs in assembly

- Compilers are much better & more patient than you are

■ But: Understanding assembly is key to machine-level execution model

- Behavior of programs in presence of bugs
 - High-level language models break down
- Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
- Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
- Creating / fighting malware

Great Reality #3: Memory Matters

■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

■ Memory performance is not uniform

- **Cache** and **virtual memory** effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```


fun(0)	-->	3.14
fun(1)	-->	3.14
fun(2)	-->	3.1399998664856
fun(3)	-->	2.00000061035156

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0) -->	3.14
fun(1) -->	3.14
fun(2) -->	3.1399998664856
fun(3) -->	2.00000061035156

Explanation:



Memory Referencing Errors

■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated

■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**

- **Must understand system to optimize performance**

- How programs compiled and executed
- How to measure program performance and identify bottlenecks
- How to improve performance without destroying code modularity and generality

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3ms


```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8ms

2.0 GHz Intel Core i7 Haswell

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

Why The Performance Differs



Great Reality #5: Computers do more than execute programs

■ They need to get data in and out

- I/O system critical to program reliability and performance

■ They communicate with each other over networks

- Many system-level issues arise in presence of network
 - Concurrent operations by multiple processes
 - Cross platform compatibility
 - Complex performance issues

Course Perspective

■ Most Systems Courses are Builder-Centric

■ Our Course is Programmer-Centric

- Purpose is to show that by knowing more about the **underlying system**
- Enable you to write programs that are more reliable and efficient

Contents

■ Part1: Course Overview

- Course Theme
- Why This Course Important? Five realities!
- Course Perspective and Policies

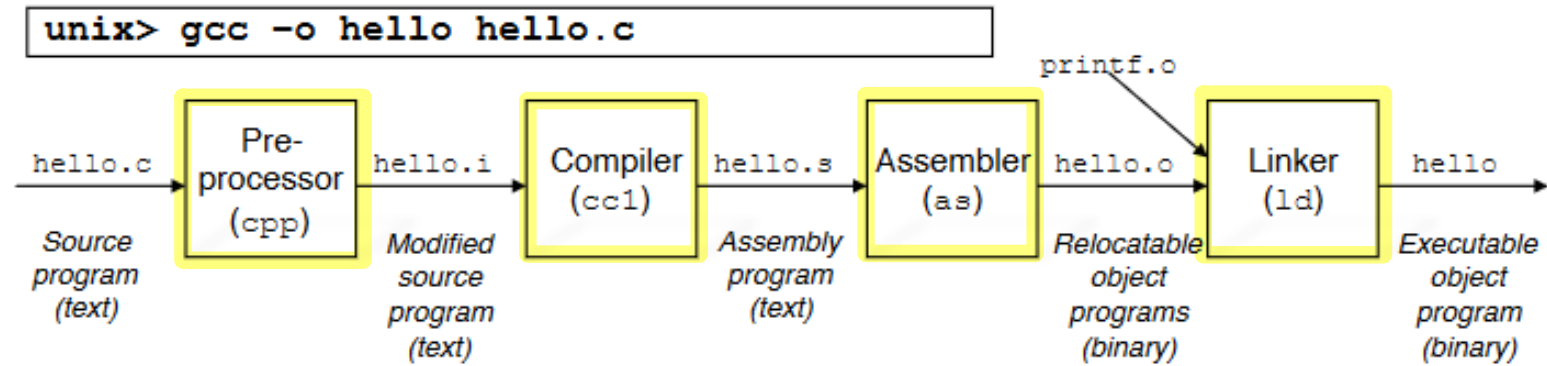
■ Part2: The Overview of Computer Systems

Hello World

- What happens and why when you run “hello” on your system?

```
/*hello world*/  
# include <stdio.h>  
  
int main()  
{  
    printf("hello, world\n");  
}
```

Programs translated by other programs



■ Pre-processing

- E.g., `#include <stdio.h>` is inserted into `hello.i`

■ Compilation (.s)

- Each statement is an assembly language program

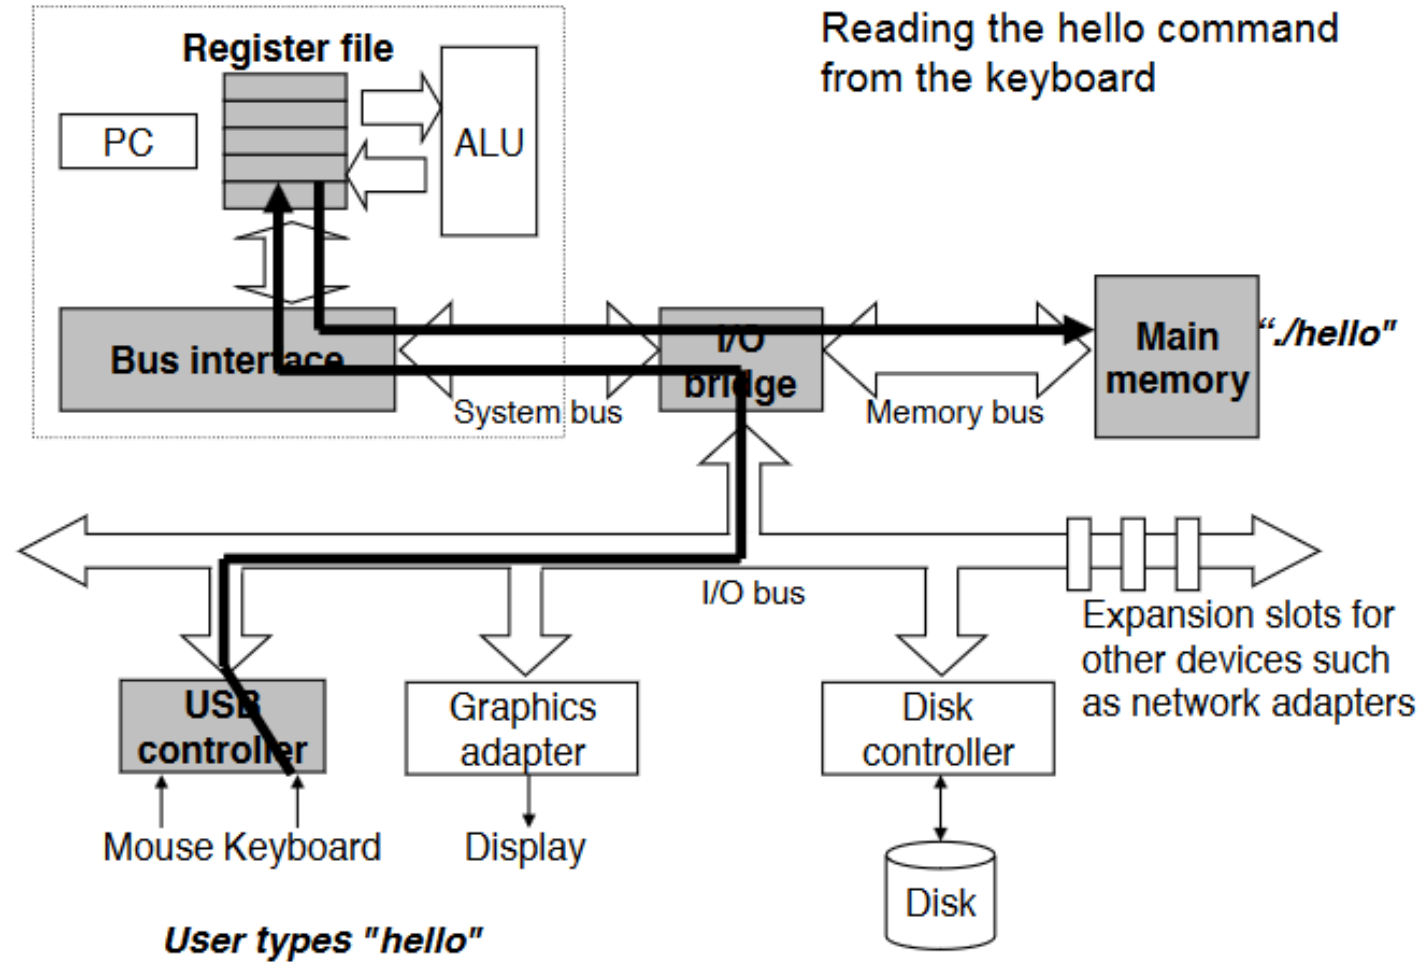
■ Assembly (.o)

- A binary file whose bytes encode machine language instructions

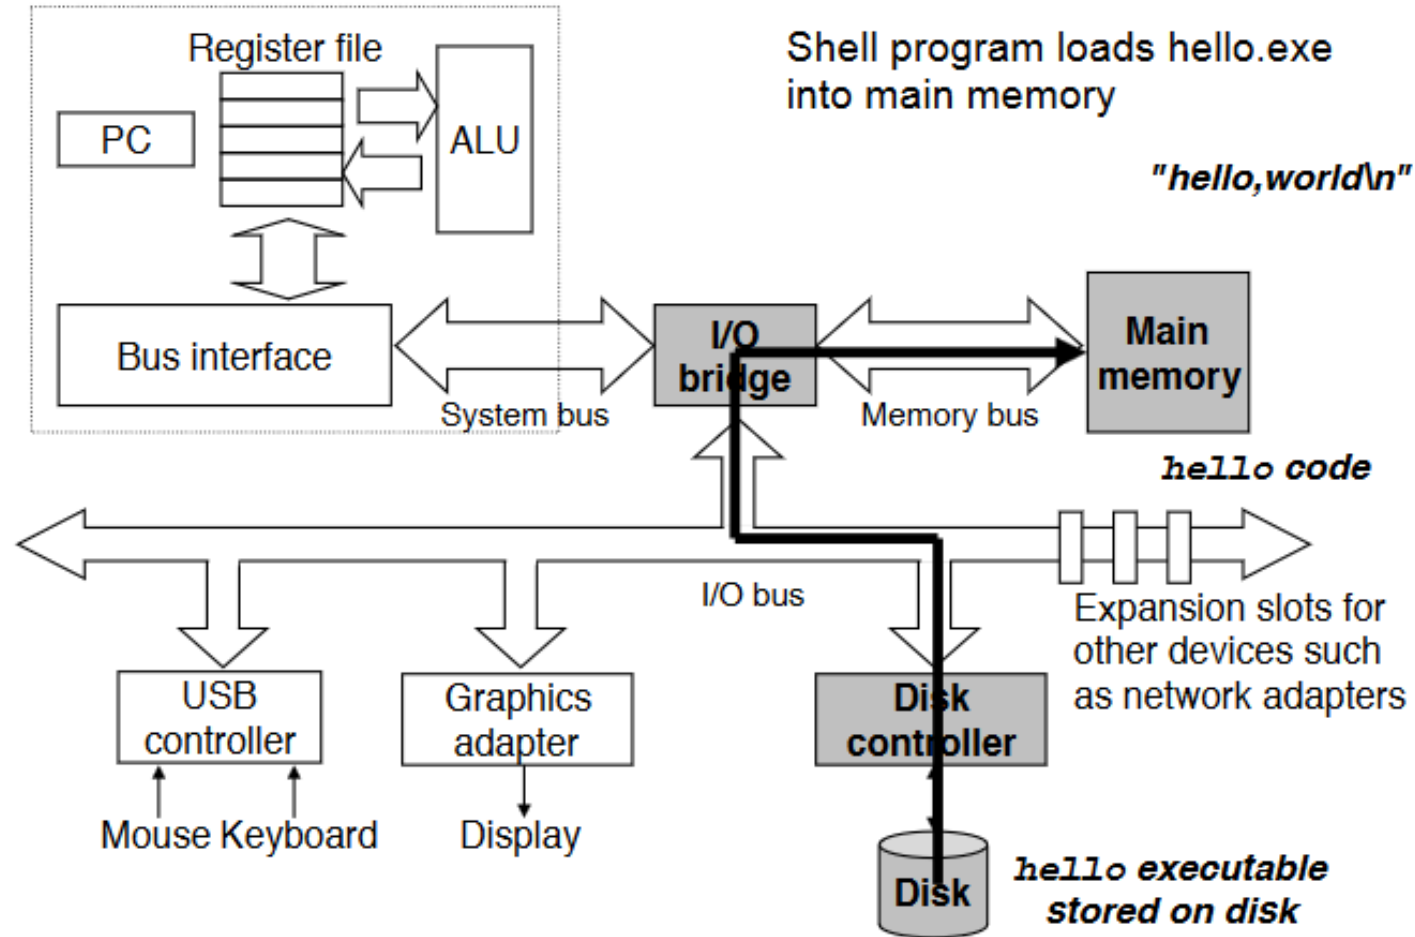
■ Linking

- Get `printf()` which resides in a separate precompiled object file

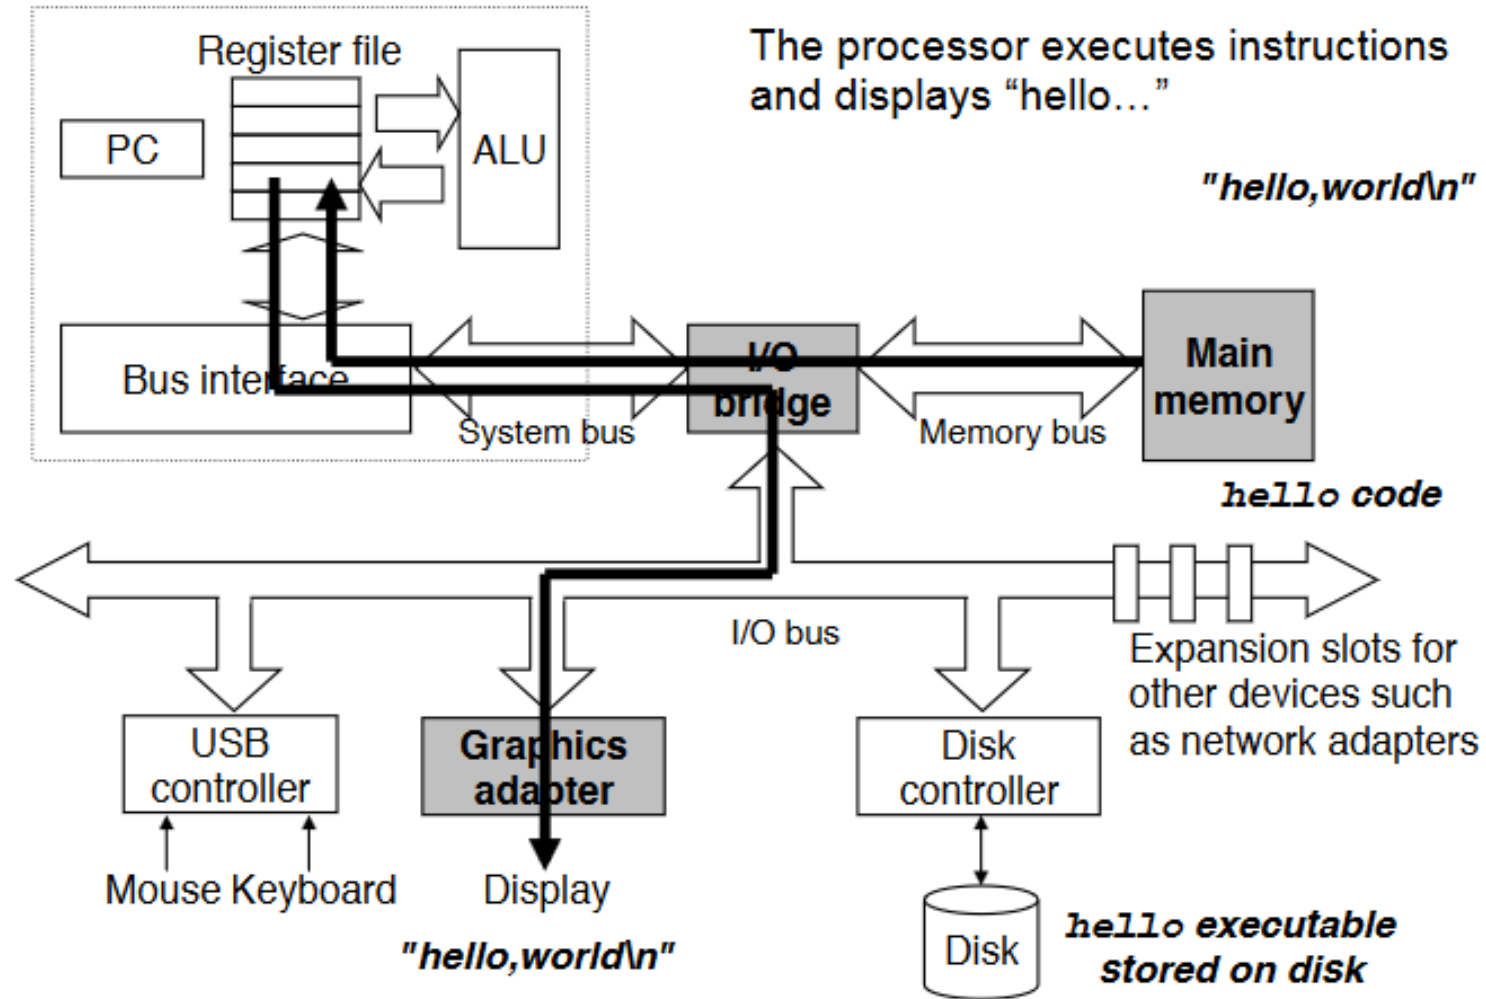
Running Hello



Running Hello



Running Hello



Running Hello

```
unix> ./hello  
hello, world  
unix>
```

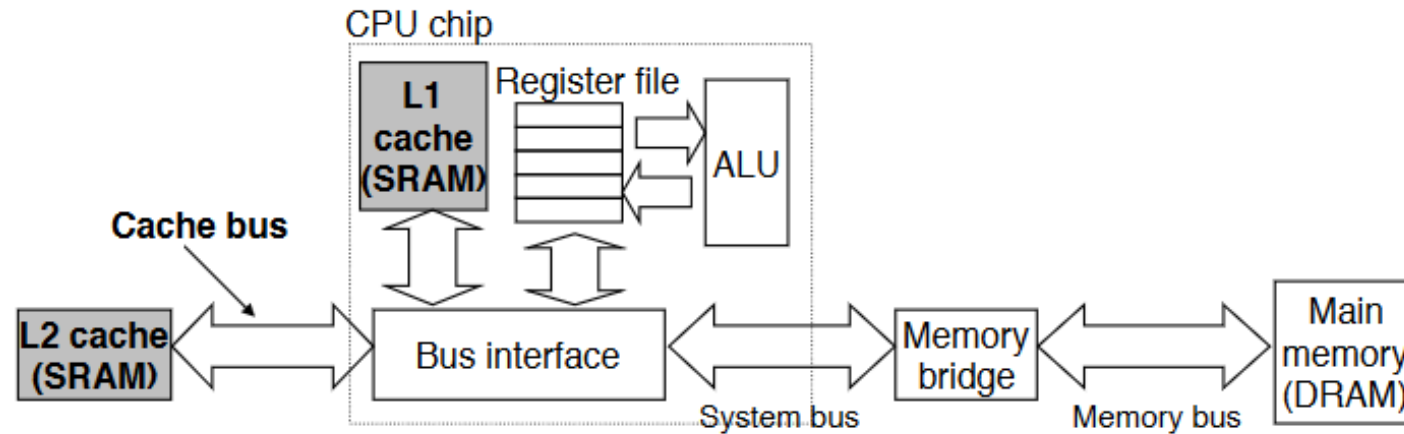
■ What's the shell?

■ What does it do?

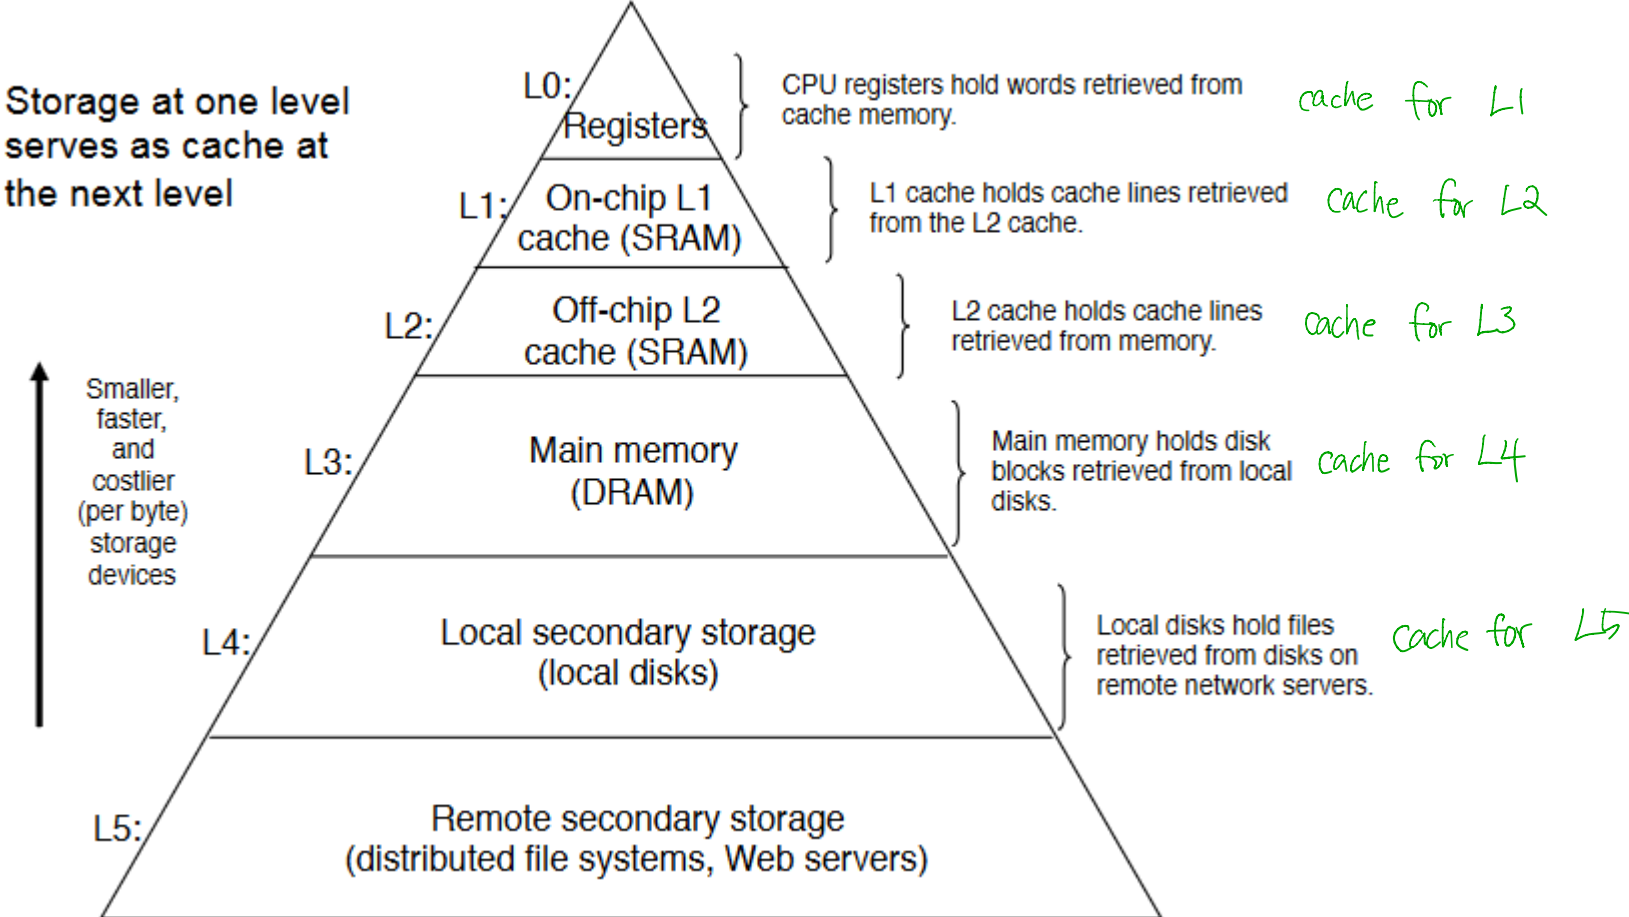
- prints a prompt
- waits for you to type command line
- loads and runs hello program
- ...

Cache Matter

- System spends a lot of time moving info. around
- Smaller storage devices are faster than larger ones
 - Register file ~ 100 Bytes & Main memory ~ millions of Bytes
- Easier and cheaper to make processors run faster than to make main memory run faster
 - Standard answer – cache




Storage Devices Form a Hierarchy



Operating System

- OS – a layer of software interposed between the application program and the hardware




- **Two primary goal**

- Provide simple and uniform mechanisms for manipulating low-level hardware devices
- Protect resources from misuse by applications

OS Abstractions

- **Files** – abstractions of I/O devices
- **Virtual memory** – abstraction for main memory and I/O devices
- **Processes** – abstractions for processor, main memory, and I/O devices



Processes

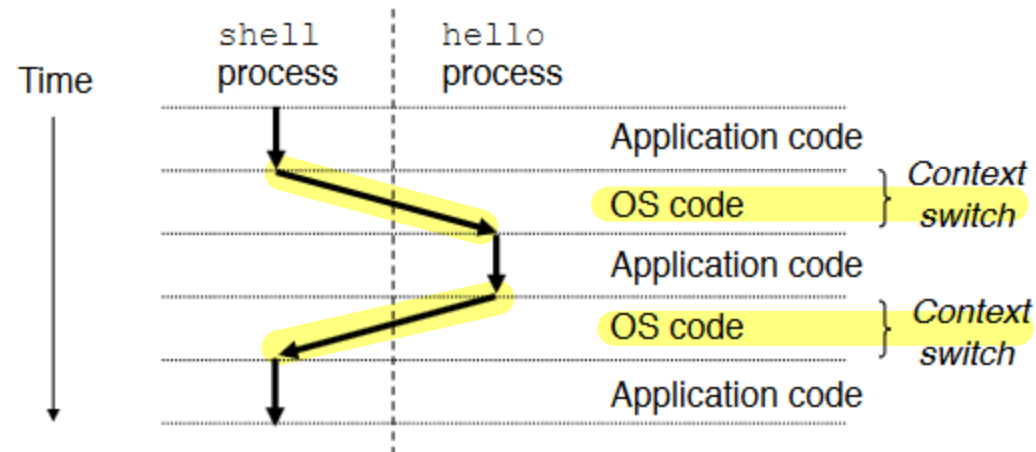
- OS provides the illusion of a dedicated machine per process

- **Process**

- OS's abstraction of a **running program**

- **Context switch**

- Saving context of one process, restoring that of another one
 - Distorted notion of time



Summary of Today's Lecture

- **A computer system is more than just hardware**

- A collection of intertwined HW & SF that must cooperate to achieve the end goal – running applications

- **The rest of the course will expand on this**

Welcome and Enjoy!