

---

# **Machine-Level Programming: Advanced Topics**

Prof. Hyuk-Yoon Kwon

<https://sites.google.com/view/seoultech-bigdata>

# Today

---

## ■ Memory Layout

## ■ Buffer Overflow

- Vulnerability
- Protection

OS → managing memory by each process

provide same address range for memory for each process

⇒ Virtual memory

Cannot access physical memory directly ⇒ OS manage physical memory &

show conceptual memory layout to the process

⇒ process can access  
the address of memory  
provided by OS & CS

# x86-64 Linux Memory Layout

... Supported by OS , all process have the same address range for the memory

## ■ Stack

- Runtime stack (8MB limit)
- E. g., local variables (return address, values ...) have been stored in registers

## ■ Heap

- Dynamically allocated as needed ... result of execution of program.
- When call malloc(), calloc(), new()

## ■ Data

determine the value of data before running the program.  $\Rightarrow$  when we compile!

- Statically allocated data
- E.g., global vars, static vars, string constants

region for data will not be changed as the program launch.

store data

## ■ Text / Shared Libraries

... related to source codes

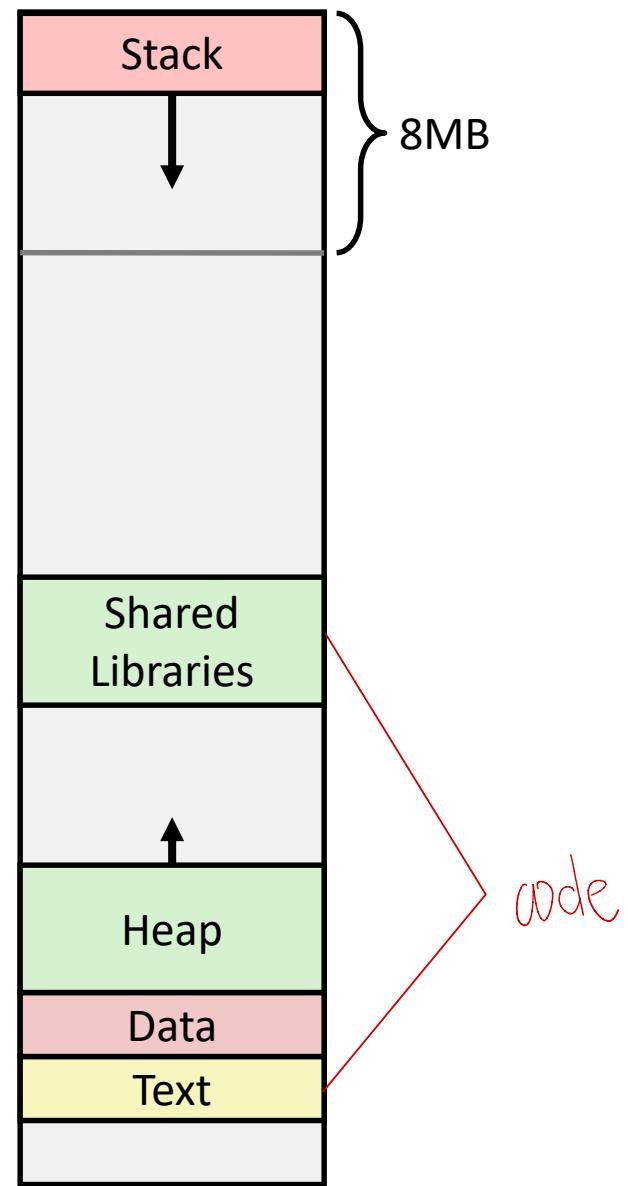
- Executable machine instructions
- Read-only

source code itself

referenced by program

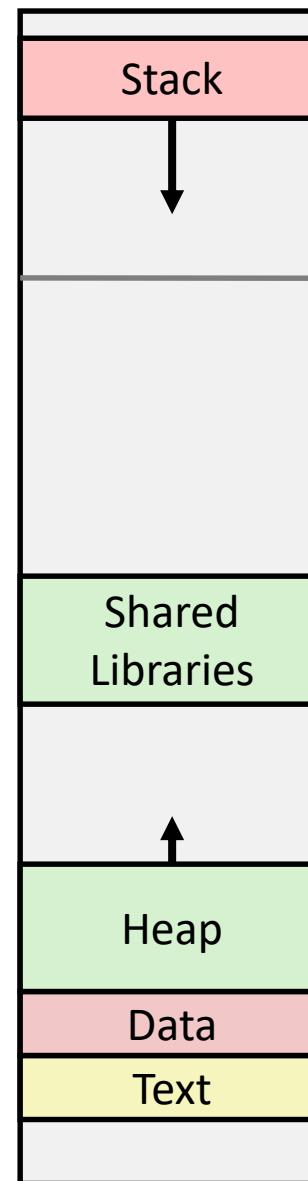
Hex Address

000000  
400000  
000000



# Memory Allocation Example

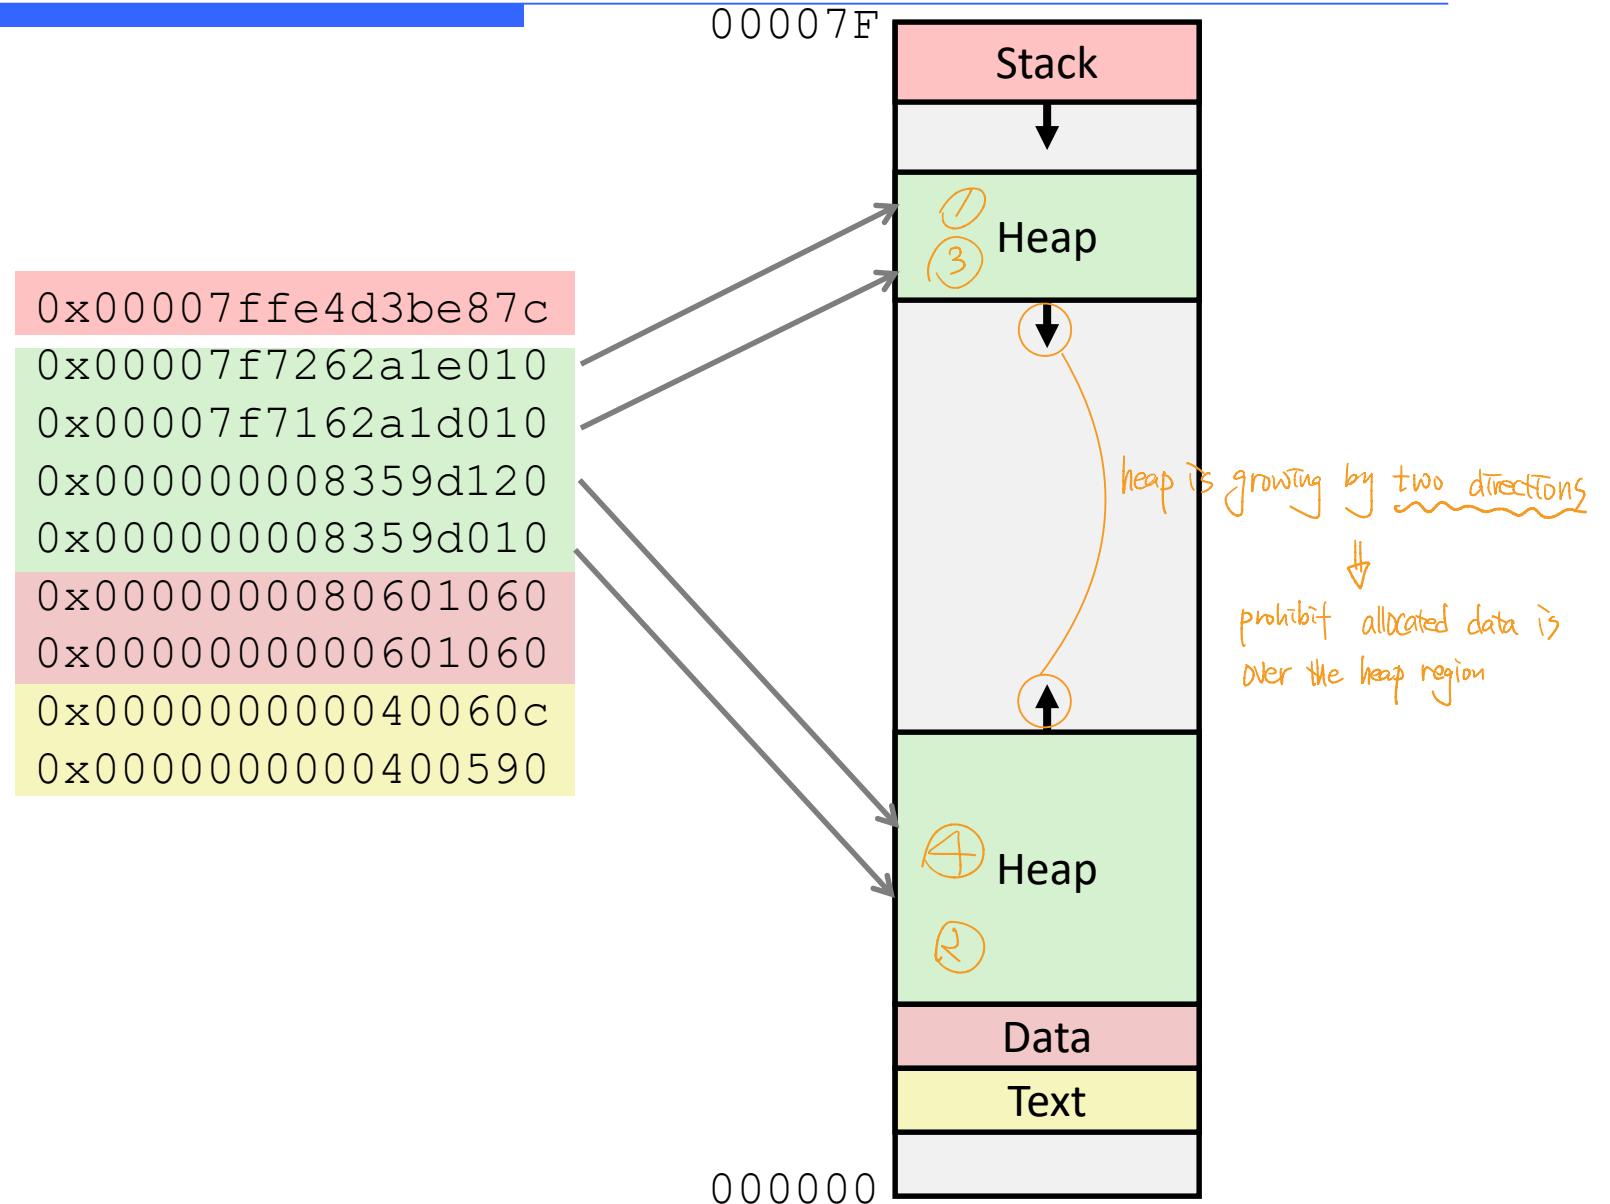
```
in [data] {  
    char big_array[1L<<24]; /* 16 MB */  
    char huge_array[1L<<31]; /* 2 GB */  
  
    int global = 0;  
  
    int useless() { return 0; }  
  
    int main ()  
    {  
        void *p1, *p2, *p3, *p4;  
        Stack { int local = 0;  
        in [text] {  
            p1 = malloc(1L << 28); /* 256 MB */  
            in [heap] { p2 = malloc(1L << 8); /* 256 B */  
            in [heap] { p3 = malloc(1L << 32); /* 4 GB */  
            in [heap] { p4 = malloc(1L << 8); /* 256 B */  
            /* Some print statements ... */  
        }  
    }  
}
```



*Where does everything go?*

# x86-64 Example Addresses

local  
p1  
p3  
p4  
p2  
big\_array  
huge\_array  
main()  
useless()



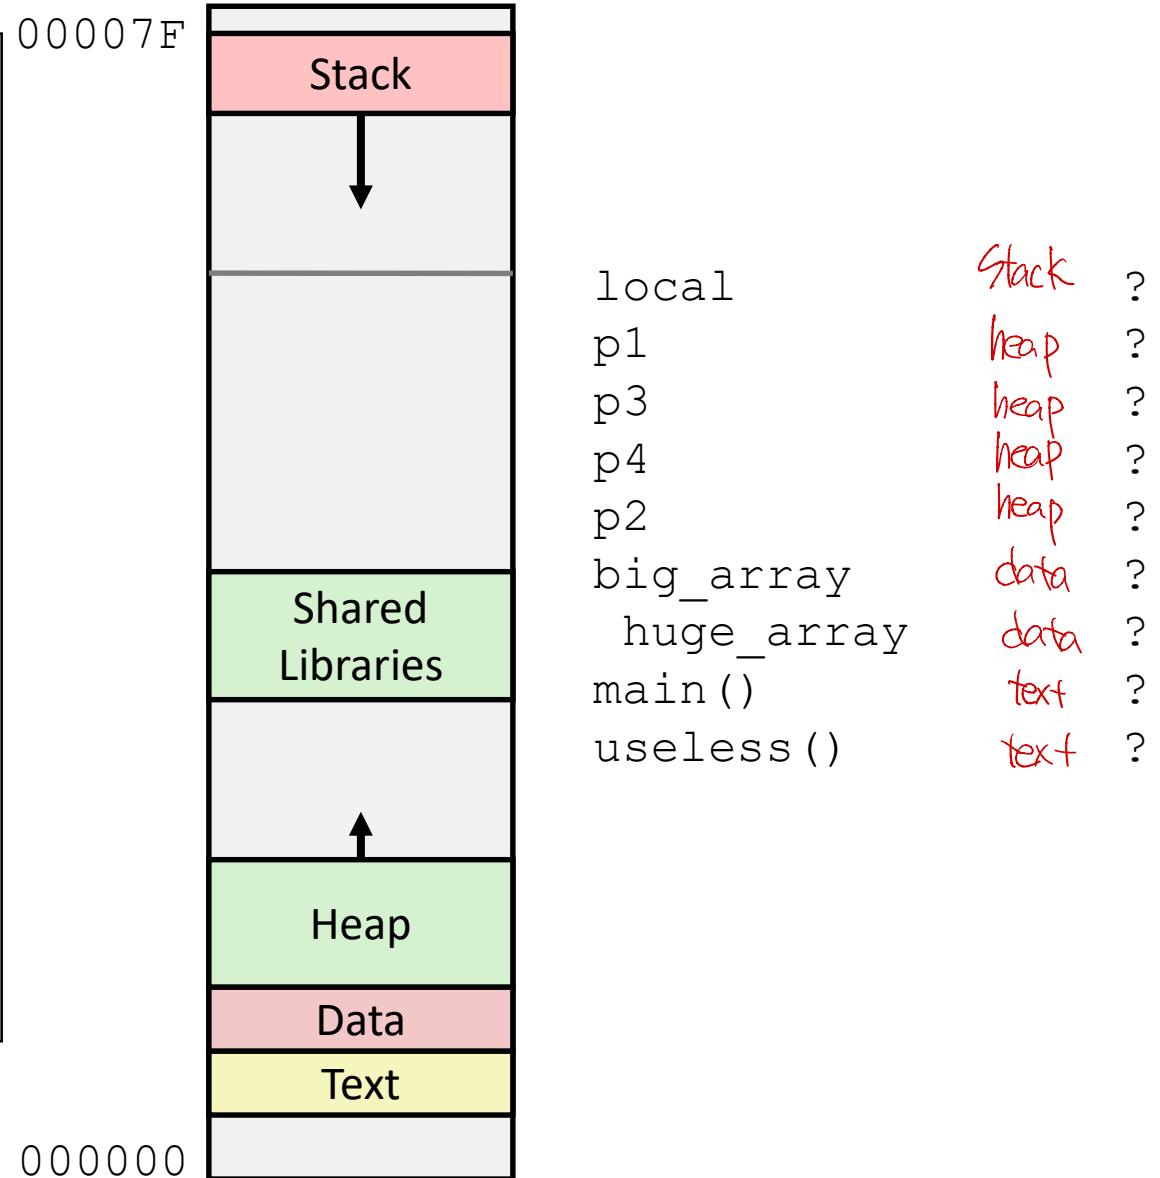
# Practice: Memory Allocation Example

```
char big_array[1L<<24]; /* 16 MB */
char huge_array[1L<<31]; /* 2 GB */

int global = 0;

int useless() { return 0; }

int main ()
{
    void *p1, *p2, *p3, *p4;
    int local = 0;
    p1 = malloc(1L << 28); /* 256 MB */
    p2 = malloc(1L << 8); /* 256 B */
    p3 = malloc(1L << 32); /* 4 GB */
    p4 = malloc(1L << 8); /* 256 B */
    /* Some print statements ... */
}
```



*Where does everything go?*

# Today

---

## ■ Memory Layout

## ■ Buffer Overflow when allocated space are overflowed.

- Vulnerability
- Protection

# Recall: Memory Referencing Bug Example

---

- Result is system specific

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

fun(0)	->	3.14
fun(1)	->	3.14
fun(2)	->	3.1399998664856
fun(3)	->	2.00000061035156
fun(4)	->	3.14
fun(6)	->	Segmentation fault

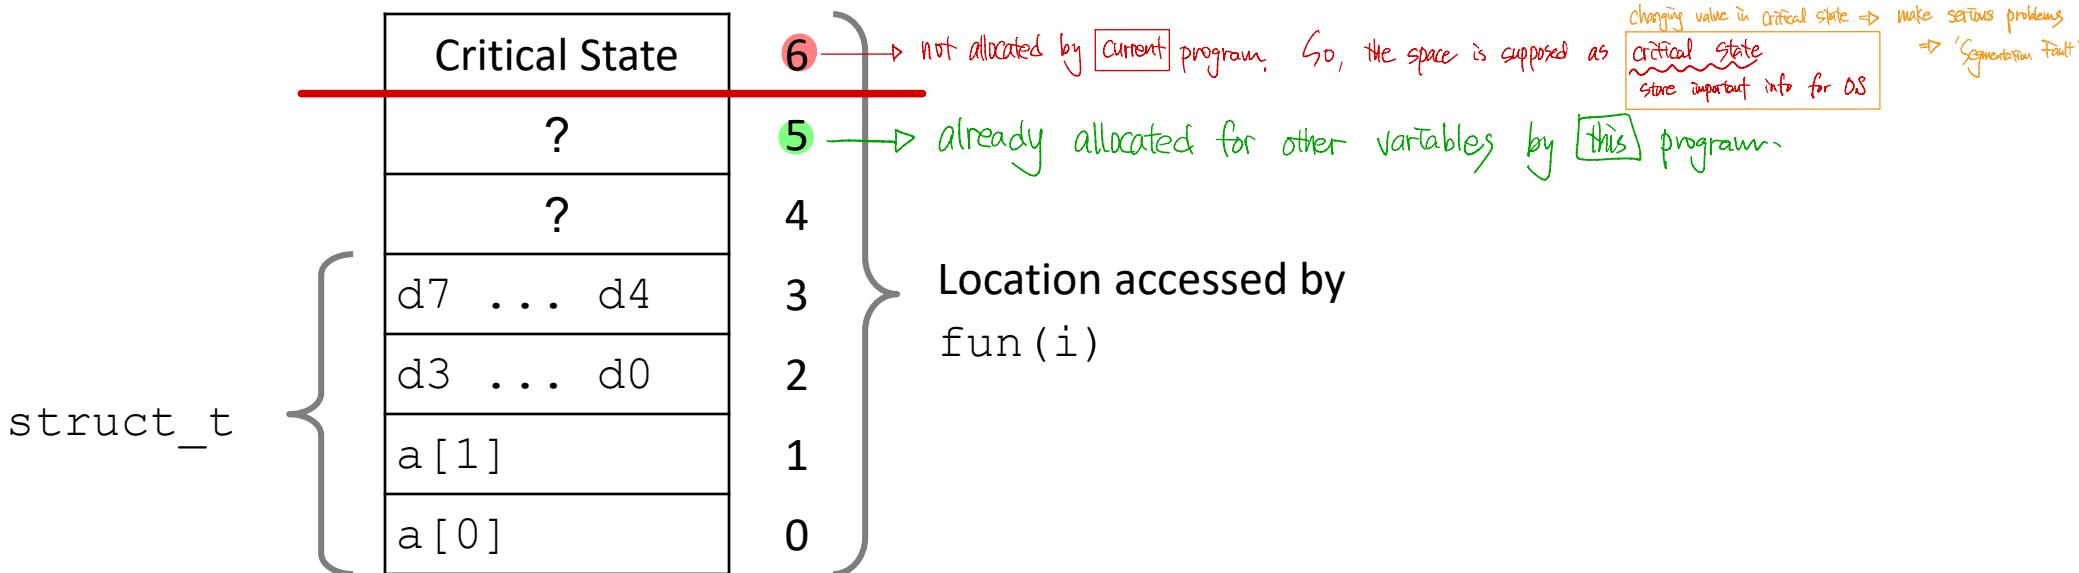
# Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

overwrite

fun(0)	->	3.14
fun(1)	->	3.14
fun(2)	->	3.1399998664856
fun(3)	->	2.00000061035156
fun(4)	->	3.14
fun(6)	->	Segmentation fault

## Explanation:



# Such problems are a BIG deal

---

## ■ Generally called a “buffer overflow”

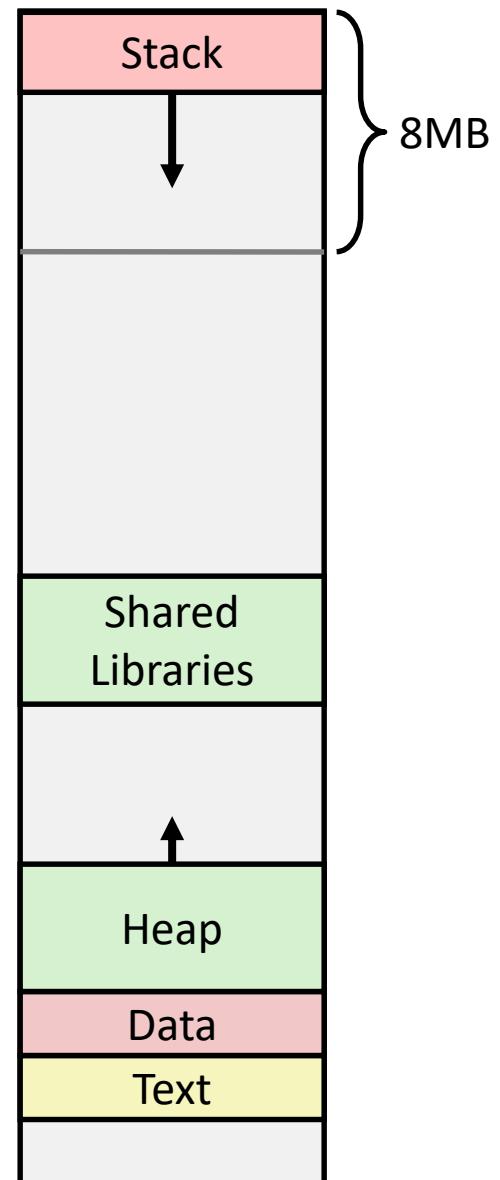
- when exceeding the memory size allocated for an array

## ■ Why a big deal?

- It's the #1 technical cause of security vulnerabilities
  - #1 overall cause is social engineering / user ignorance

## ■ Most common form

- Unchecked lengths on string inputs
- Particularly for bounded character arrays on the stack
  - sometimes referred to as stack smashing



# String Library Code

## Implementation of Unix function gets ()

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

can occur buffer overflow!

If input string exceed the size of allocated memory  
⇒ Store the string into other region in the memory  
OVERRWRITE! ... security vulnerability

- No way to specify limit on number of characters to read

## Similar problems with other library functions

- strcpy, strcat:** Copy strings of arbitrary length
- scanf, fscanf, sscanf,** when given %s conversion specification

# Vulnerable Buffer Code

---

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
void call_echo() {
    echo();
}
```

# Buffer Overflow Disassembly

## echo:

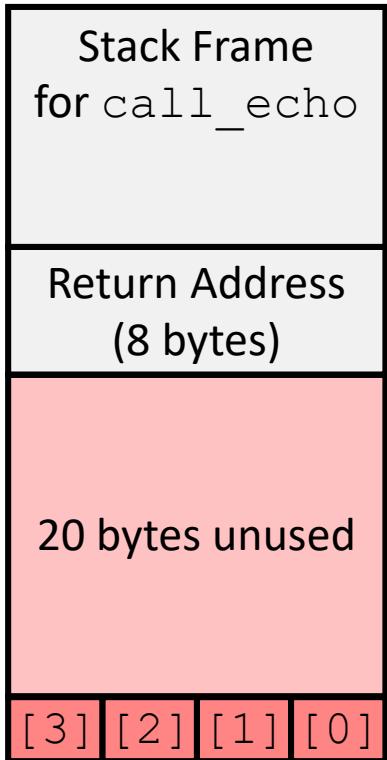
```
00000000004006cf <echo>:  
4006cf: 48 83 ec 18          sub    $0x18,%rsp  
4006d3: 48 89 e7          mov    %rsp,%rdi  
4006d6: e8 a5 ff ff ff      callq   400680 <gets>  
4006db: 48 89 e7          mov    %rsp,%rdi  
4006de: e8 3d fe ff ff      callq   400520 <puts@plt>  
4006e3: 48 83 c4 18          add    $0x18,%rsp  
4006e7: c3                  retq
```

## call\_echo:

```
4006e8: 48 83 ec 08          sub    $0x8,%rsp  
4006ec: b8 00 00 00 00      mov    $0x0,%eax  
4006f1: e8 d9 ff ff ff      callq   4006cf <echo>  
4006f6: 48 83 c4 08          add    $0x8,%rsp  
4006fa: c3                  retq
```

# Buffer Overflow Stack

*Before call to gets*

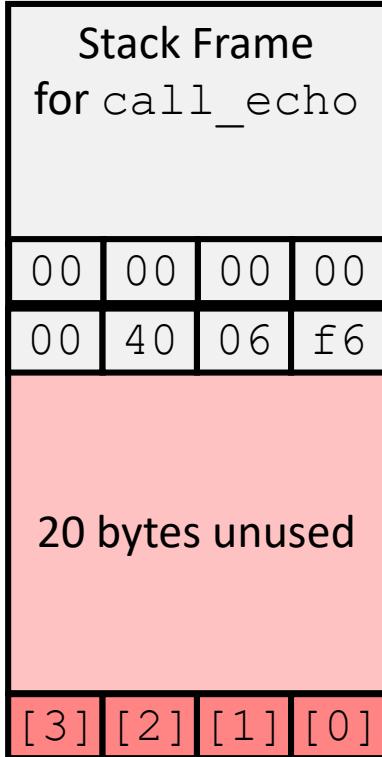


```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
    subq $24, %rsp
    movq %rsp, %rdi
    call gets
    . . .
```

# Buffer Overflow Stack Example

*Before call to gets*



```
void echo ()  
{  
    char buf[4];  
    gets(buf);  
    . . .  
}
```

```
echo:  
    subq $24, %rsp  
    movq %rsp, %rdi  
    call gets  
    . . .
```

call\_echo:

```
. . .  
4006f1: callq 4006cf <echo>  
4006f6: add $0x8,%rsp  
. . .
```

# Buffer Overflow Stack Example #1

*After call to gets*

Stack Frame for call_echo			
00	00	00	00
00	40	06	f6
00	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

```
void echo()
{
    char buf[4];
    gets(buf);
    ...
}
```

```
echo:
    subq $24, %rsp
    movq %rsp, %rdi
    call gets
    ...
}
```

**call\_echo:**

```
...
4006f1: callq 4006cf <echo>
4006f6: add    $0x8,%rsp
...
```

buf ← %rsp

```
unix>./bufdemo-nsp
Type a string:01234567890123456789012
01234567890123456789012
```

Overflowed buffer, but did not corrupt state

24 bytes

including '\0'

# Buffer Overflow Stack Example #2

*After call to gets*

Stack Frame for call_echo			
00	00	00	00
00	40	00	34
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

```
void echo()
{
    char buf[4];
    gets(buf);
    ...
}
```

```
echo:
    subq $24, %rsp
    movq %rsp, %rdi
    call gets
    ...
}
```

**call\_echo:**

```
...
4006f1: callq 4006cf <echo>
4006f6: add    $0x8,%rsp
...
```

buf ← %rsp

```
unix>./bufdemo-nsp
Type a string:0123456789012345678901234
Segmentation Fault
```

Overflowed buffer and corrupted return pointer

26 bytes  
including '\0'

# Buffer Overflow Stack Example #3

*After call to gets*

Stack Frame for call_echo			
00	00	00	00
00	40	06	00
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

```
void echo()
{
    char buf[4];
    gets(buf);
    ...
}
```

```
echo:
    subq $24, %rsp
    movq %rsp, %rdi
    call gets
    ...
}
```

**call\_echo:**

```
...
4006f1: callq 4006cf <echo>
4006f6: add    $0x8,%rsp
...
```

buf ← %rsp

```
unix>./bufdemo-nsp
Type a string:012345678901234567890123
012345678901234567890123
```

25 bytes  
including '\0'  
↗ just luck

Overflowed buffer, corrupted return pointer, but program seems to work!

# Buffer Overflow Stack Example #3 Explained

*After call to gets*

Stack Frame for call_echo			
00	00	00	00
00	40	06	00
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

register\_tm\_clones:

```
...  
400600: mov    %rsp, %rbp  
400603: mov    %rax, %rdx  
400606: shr    $0x3f, %rdx  
40060a: add    %rdx, %rax  
40060d: sar    %rax  
400610: jne    400614  
400612: pop    %rbp  
400613: retq
```

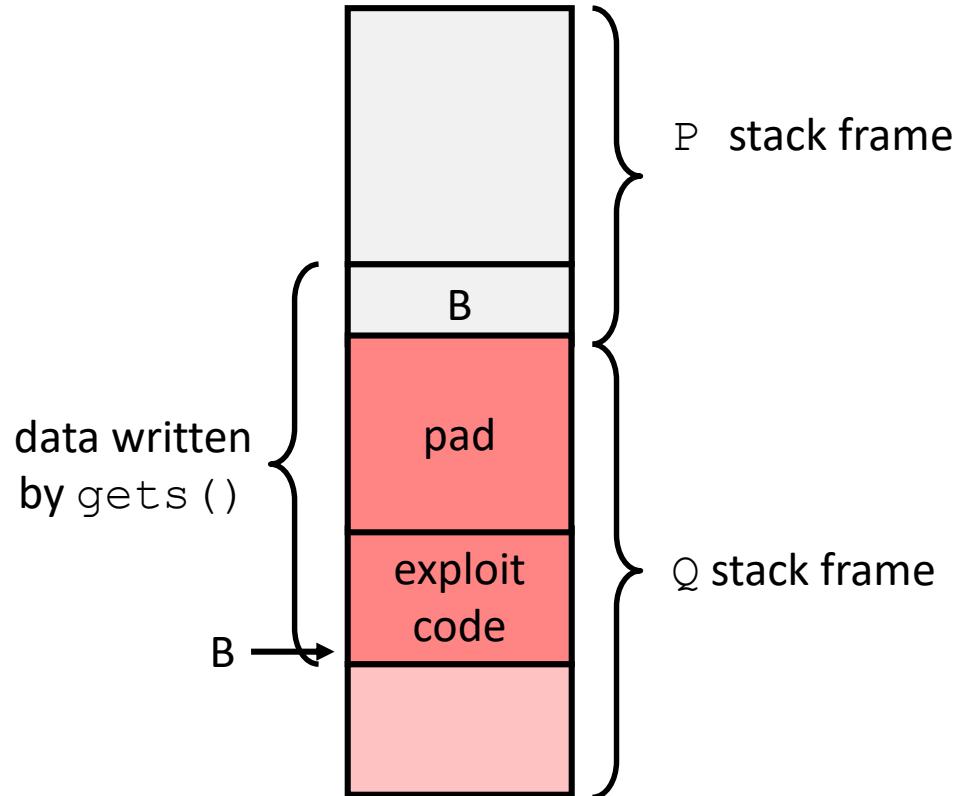
buf ← %rsp

“Returns” to unrelated code  
Lots of things happen, without modifying critical state  
Eventually executes `retq` back to `main`

# Code Injection Attacks

```
void P() {  
    Q();  
    ...  
}  
  
int Q() {  
    char buf[64];  
    gets(buf);  
    ...  
    return ...;  
}
```

Stack after call to gets ()



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer B
- When Q executes `ret`, will jump to exploit code

# Practice: Memory Referencing Bug Example

---

- Result is system specific

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

fun(0)	->	3.14
fun(1)	->	3.14
fun(2)	->	3.1399998664856
fun(3)	->	2.00000061035156
fun(4)	->	3.14
fun(6)	->	Segmentation fault

# Exploits Based on Buffer Overflows

---

■ ***Buffer overflow bugs can allow remote machines to execute arbitrary code on victim machines***

■ Examples across the decades

- Original “Internet worm” (1988)
- “IM wars” (1999)
- Twilight hack on Wii (2000s)
- ... and many, many more

# Example: the original Internet worm (1988)

---

## ■ Exploited a few vulnerabilities to spread

- Early versions of the finger server (fingerd) used `gets()` to read the argument sent by the client:
  - `finger droh@cs.cmu.edu`
- Worm attacked fingerd server by sending phony argument:
  - `finger "exploit-code padding new-return-address"`
  - exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.

## ■ Once on a machine, scanned for other machines to attack

- invaded ~6000 computers in hours (10% of the Internet ☺)
  - see June 1989 article in *Comm. of the ACM*
- the young author of the worm was prosecuted...

# OK, what to do about buffer overflow attacks

---

- Avoid overflow vulnerabilities
- Employ system-level protections
- Have compiler use “stack canaries”
- Lets talk about each...

# 1. Avoid Overflow Vulnerabilities in Code (!)

---

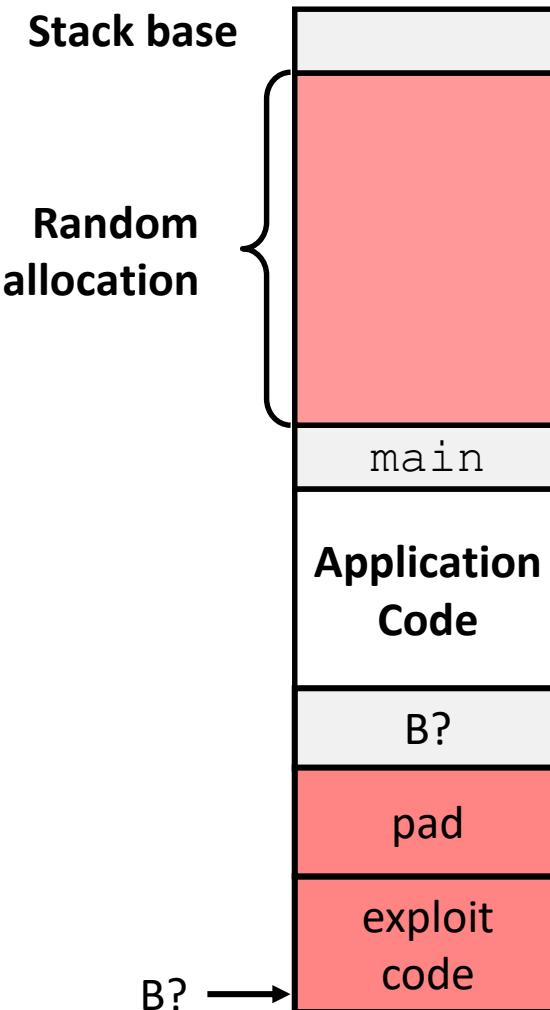
```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    fgets(buf, 4, stdin);
    puts(buf);
}
```

- For example, use library routines that limit string lengths
  - **fgets** instead of **gets**
  - **strncpy** instead of **strcpy**

## 2. System-Level Protections can help

### ■ Randomized stack offsets

- At start of program, allocate random amount of space on stack
- Shifts stack addresses for entire program
- Makes it difficult for hacker to predict beginning of inserted code

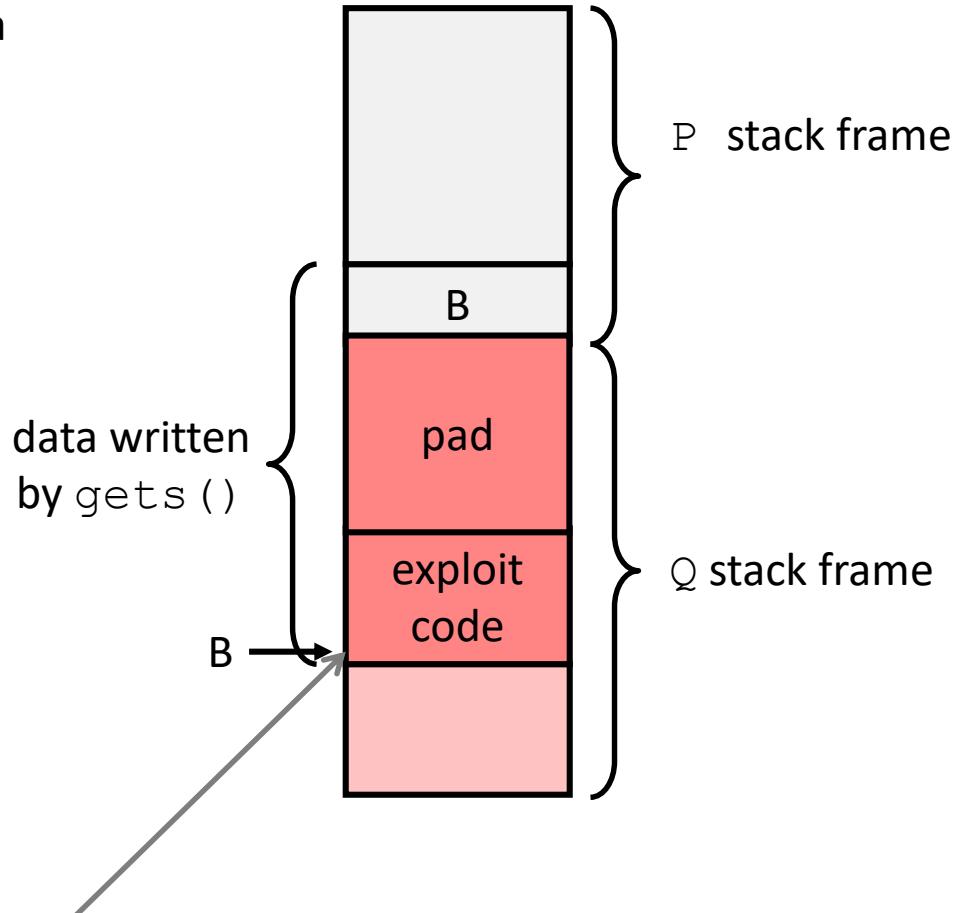


## 2. System-Level Protections can help

### ■ Nonexecutable code segments

- In traditional x86, can mark region of memory as either “read-only” or “writeable”
  - Can execute anything readable
- X86-64 added explicit “execute” permission
- Stack marked as non-executable

Stack after call to gets ()



Any attempt to execute this code will fail

### 3. Stack Canaries can help

---

#### ■ Idea

- Place special value (“canary”) on stack just beyond buffer
- Check for corruption before exiting function

#### ■ GCC Implementation

- **-fstack-protector**
- Now the default (disabled earlier)

```
unix>./bufdemo-sp
Type a string:0123456
0123456
```

```
unix>./bufdemo-sp
Type a string:01234567
*** stack smashing detected ***
```

# Protected Buffer Disassembly

using canary

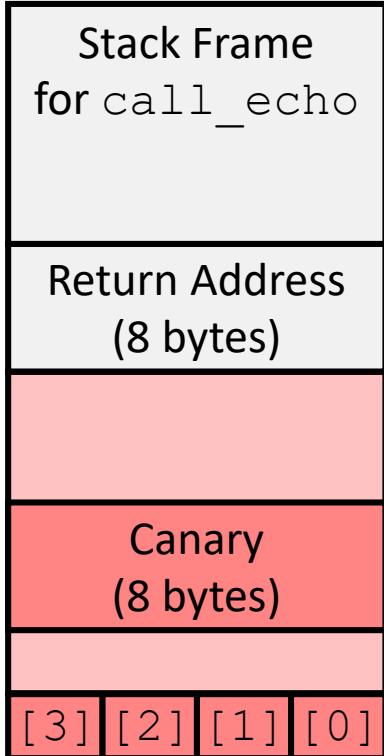
echo:

```
40072f: sub    $0x18,%rsp
400733: mov    %fs:0x28,%rax
40073c: mov    %rax,0x8(%rsp)
400741: xor    %eax,%eax
400743: mov    %rsp,%rdi
400746: callq  4006e0 <gets>
40074b: mov    %rsp,%rdi
40074e: callq  400570 <puts@plt>
400753: mov    0x8(%rsp),%rax
400758: xor    %fs:0x28,%rax
400761: je     400768 <echo+0x39>
400763: callq  400580 <__stack_chk_fail@plt>
400768: add    $0x18,%rsp
40076c: retq
```

compare canary value  
before & after executing function.

# Setting Up Canary

*Before call to gets*



```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

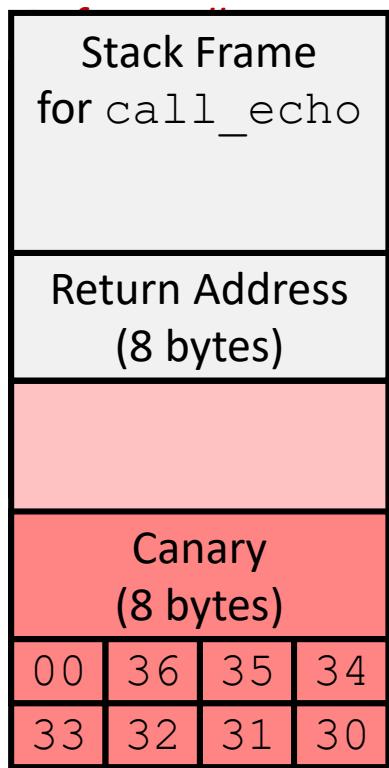
→ middle between return address & actual value in stack  
if buffer overflow & overwriting occur, canary must be changed  
Checking canary before return to origin function!  
buf ← %rsp

```
echo:
```

```
    . . .
    movq    %fs:40, %rax # Get canary
    movq    %rax, 8(%rsp) # Place on stack
    xorl    %eax, %eax # Erase canary
    . . .
```

# Checking Canary

*After call to gets*



```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

Input: 0123456

```
echo:
    . . .
    movq    8(%rsp), %rax      # Retrieve from
stack
    xorq    %fs:40, %rax      # Compare to canary
    je     .L6                  # If same, OK
    call   __stack_chk_fail    # FAIL
.L6:   . . .
```

# Return-Oriented Programming Attacks

## ■ Challenge (for hackers)

- Stack randomization makes it hard to predict buffer location
- Marking stack nonexecutable makes it hard to insert binary code

## ■ Alternative Strategy

- Use existing code + shared libraries ) without writing any codes
  - E.g., library code from stdlib
- String together fragments to achieve overall desired outcome multiple separate steps to execute codes
- *Does not overcome stack canaries*

## ■ Construct program from gadgets code blocks

- Sequence of instructions ending in `ret`
  - Encoded by single byte `0xc3`
- Code positions fixed from run to run
- Code is executable

# Gadget Example

```
long ab_plus_c  
    (long a, long b, long c)  
{  
    return a*b + c;  
}
```

```
00000000004004d0 <ab_plus_c>:  
4004d0: 48 0f af fe  imul %rsi,%rdi  
4004d4: 48 8d 04 17  lea (%rdi,%rdx,1),%rax  
4004d8: c3                retq
```



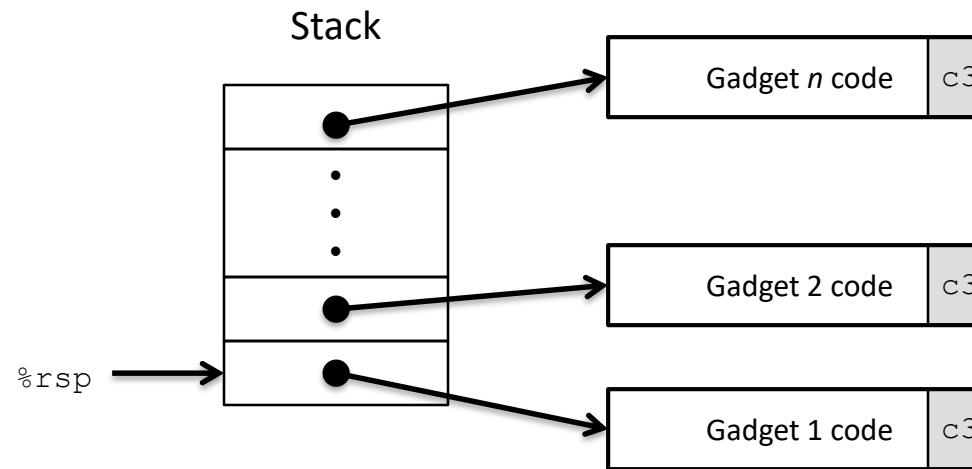
$\text{rax} \leftarrow \text{rdi} + \text{rdx}$

Gadget address = 0x4004d4

- Use tail end of existing functions

# ROP Execution

---



- Trigger with `ret` instruction
  - Will start executing Gadget 1
  
- Final `ret` in each gadget will start next one

# Today

---

- Storage technologies and trends
- Locality of reference
- Caching in the memory hierarchy

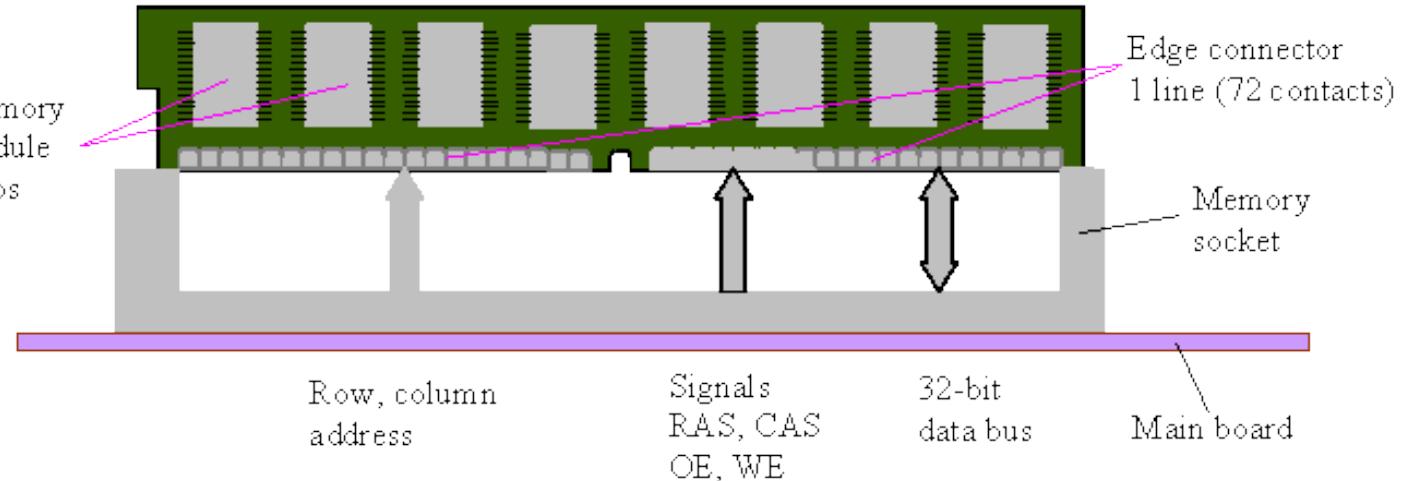
# Random-Access Memory (RAM)

## ■ Key features

- RAM is traditionally packaged as a chip.  
*Smallest unit for RAM*
- Basic storage unit is normally a cell (one bit per cell).
- Multiple RAM chips form a memory.

## ■ RAM comes in two varieties:

- SRAM (Static RAM)
- DRAM (Dynamic RAM)



# SRAM vs DRAM Summary

Trans. per bit	Access time	Needs refresh?	Cost	Applications
SRAM    4 or 6	1X	No	100x	Cache memories
DRAM    1	10X	Yes	1X	Main memories, frame buffers

# Nonvolatile Memories

## ■ DRAM and SRAM are volatile memories

- Lose information if powered off.

like disk!

## ■ Nonvolatile memories retain value even if powered off

- Read-only memory (**ROM**): programmed during production
- Programmable ROM (**PROM**): can be programmed once
- Eraseable PROM (**EPROM**): can be bulk erased (UV, X-Ray)
- Electrically eraseable PROM (**EEPROM**): electronic erase capability
- Flash memory: EEPROMs. with partial (block-level) erase capability
  - Wears out after about 100,000 erasings

entire erasing

못쓰게 되다.

## ■ Uses for Nonvolatile Memories

- Firmware programs stored in a ROM (BIOS, controllers for disks, network cards, graphics accelerators, security subsystems,...)  
what is it?
- Solid state disks (replace rotating disks in thumb drives, smart phones, mp3 players, tablets, laptops,...)  
SSD
- Disk caches

Minimizing gap between memory & disk

⇒ region called disk cache in the disk

store frequently used data ⇒ improving performance of accessing disk



PC boot ~ OS boot : check status of HW, change configuration

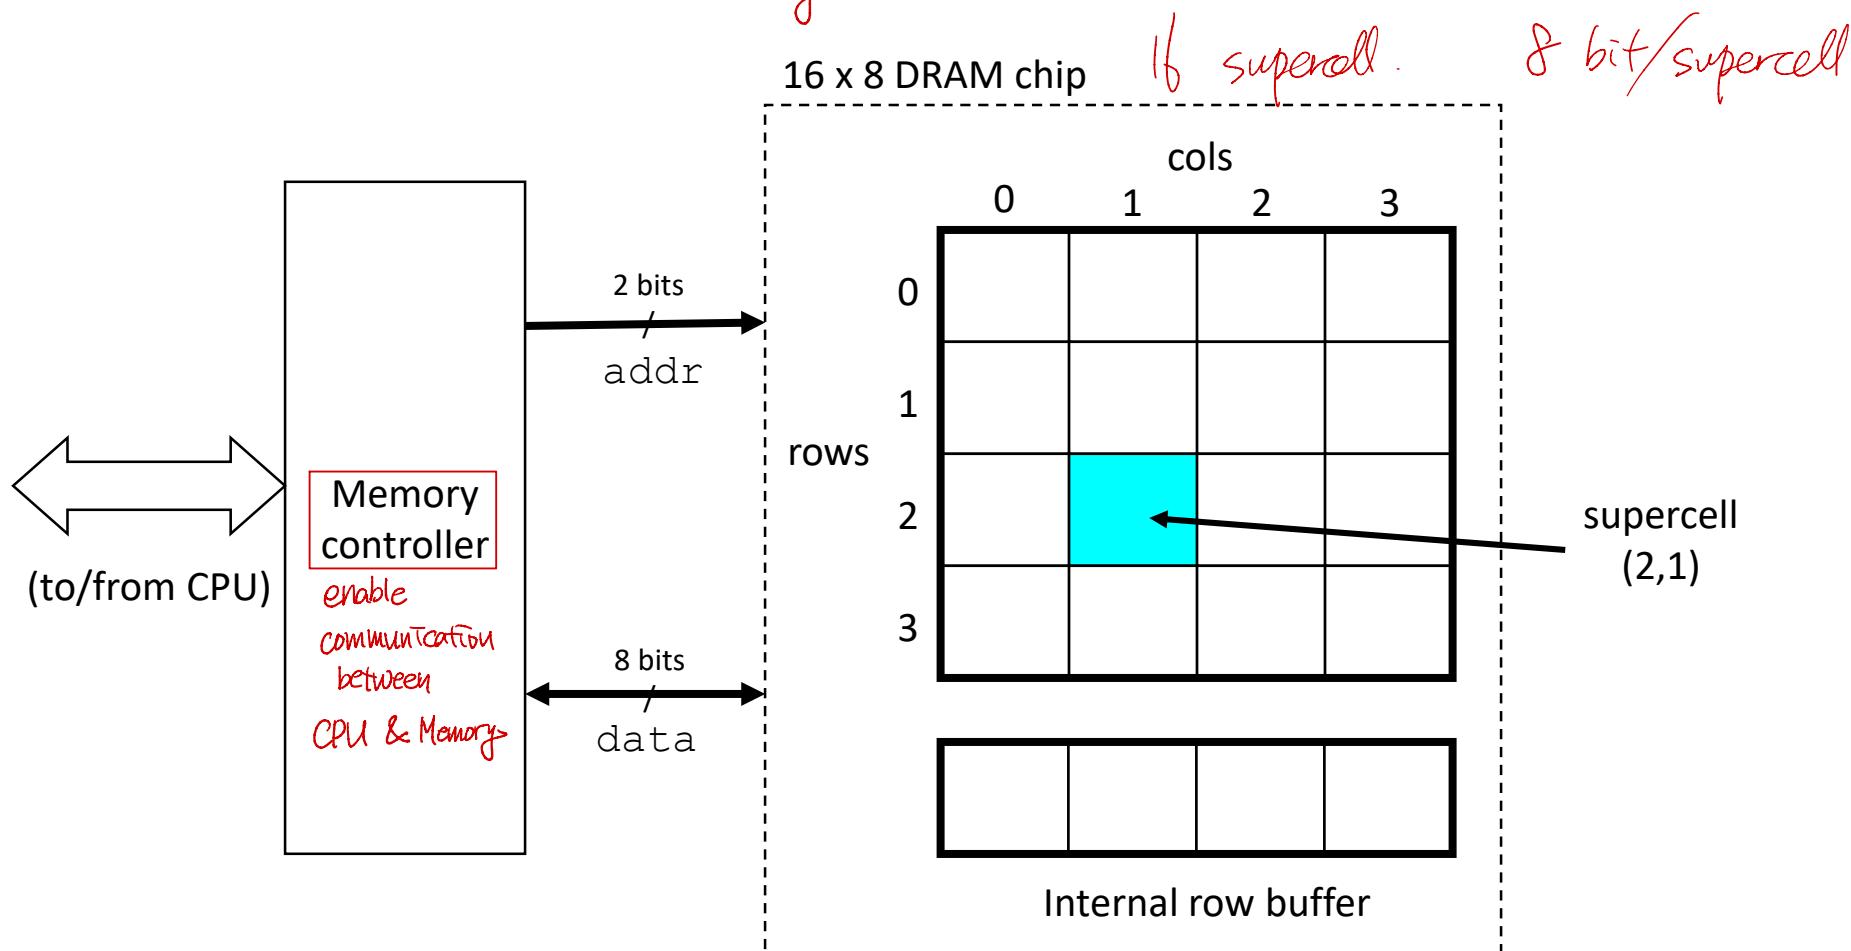
even if OS has trouble, BIOS/UEFI is always working because they are stored in non-volatile memories!

# Conventional DRAM Organization

## ■ $d \times w$ DRAM:

- dw total bits organized as d supercells of size w bits

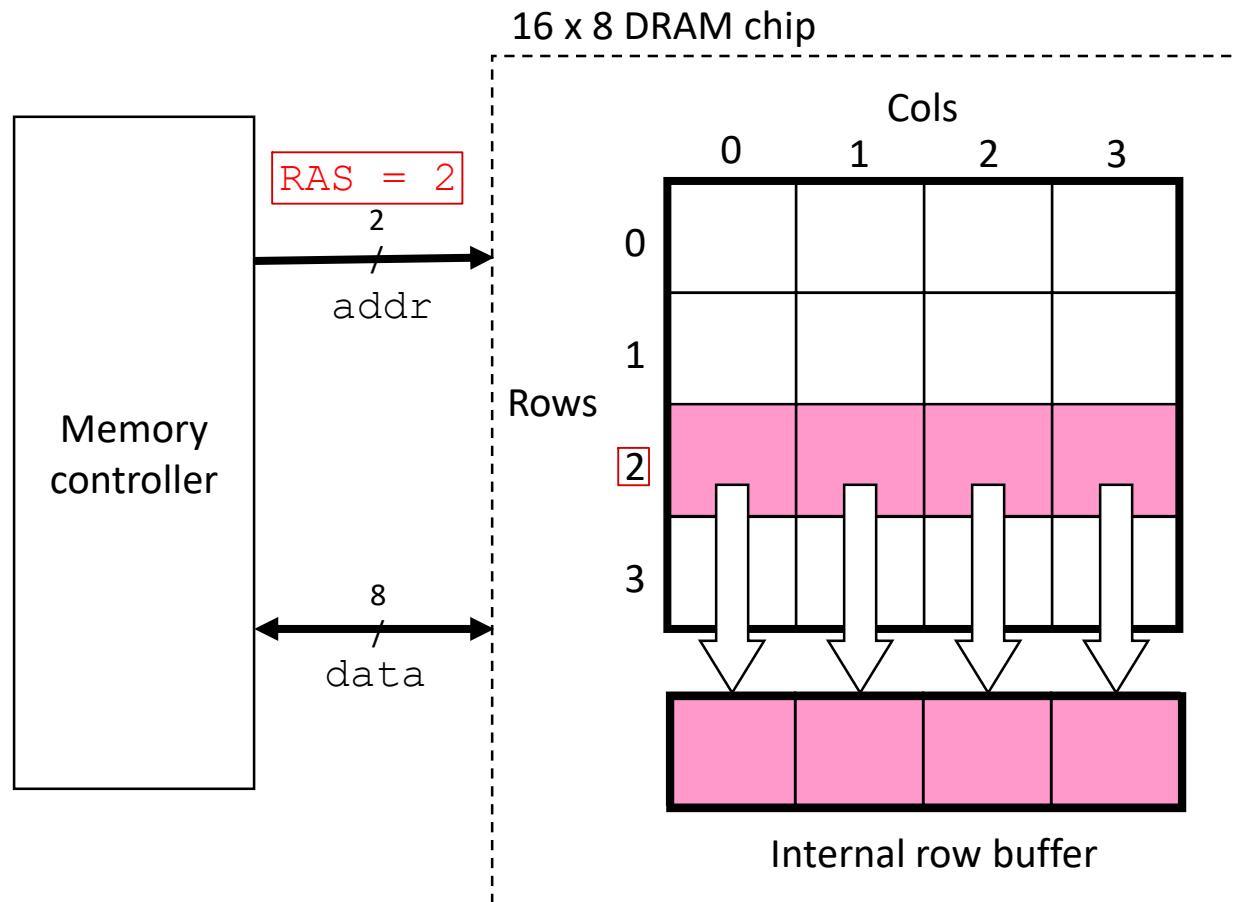
each element of 2D array model of DRAM



# Reading DRAM Supercell (2,1)

Step 1(a): Row access strobe (**RAS**) selects row 2.

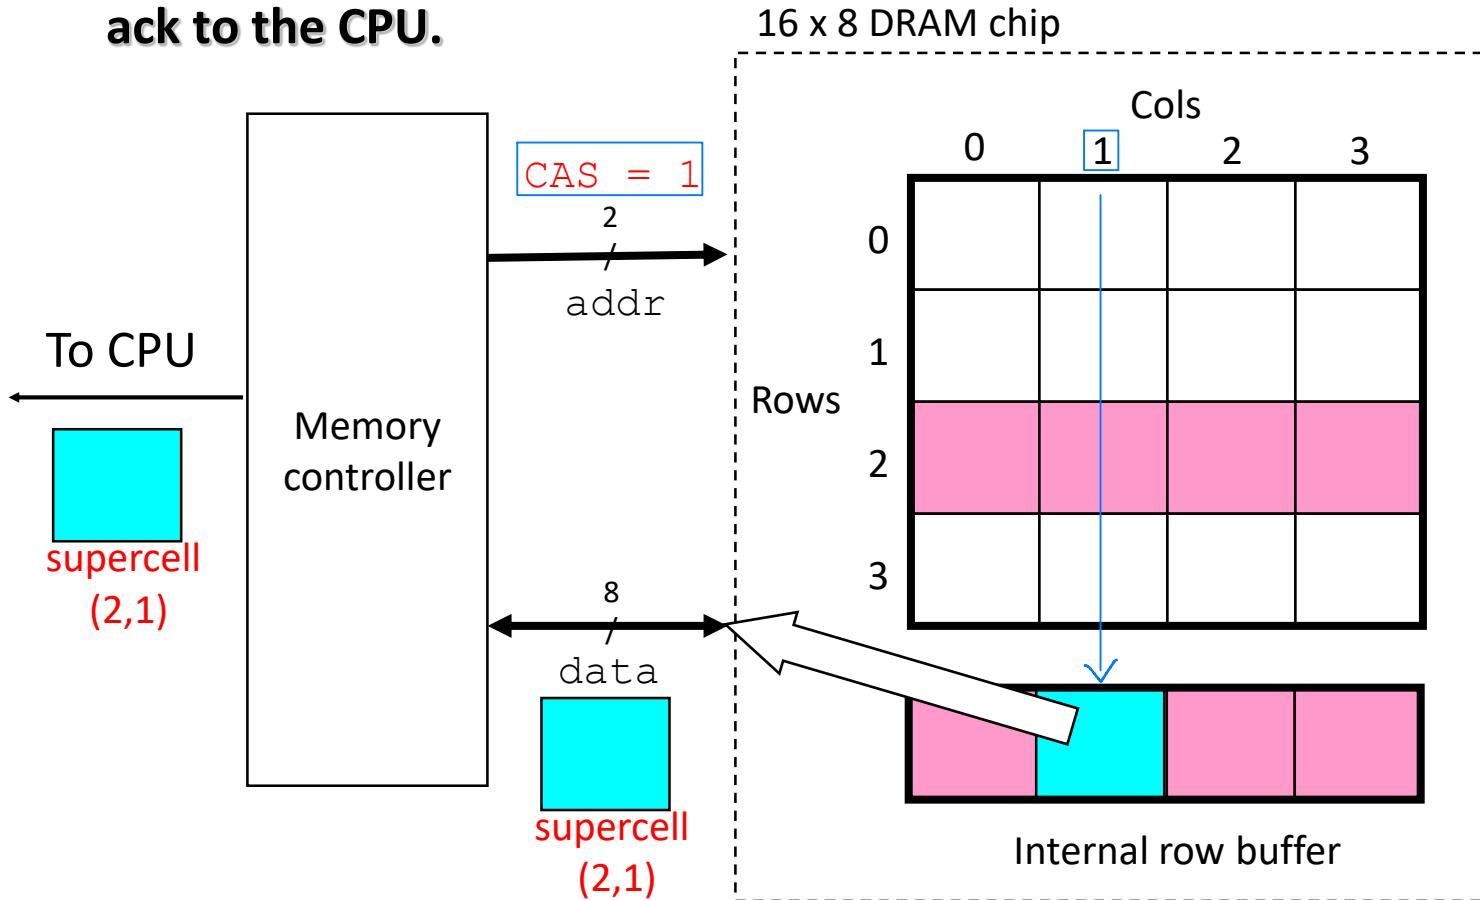
Step 1(b): Row 2 copied from DRAM array to row buffer.



# Reading DRAM Supercell (2,1)

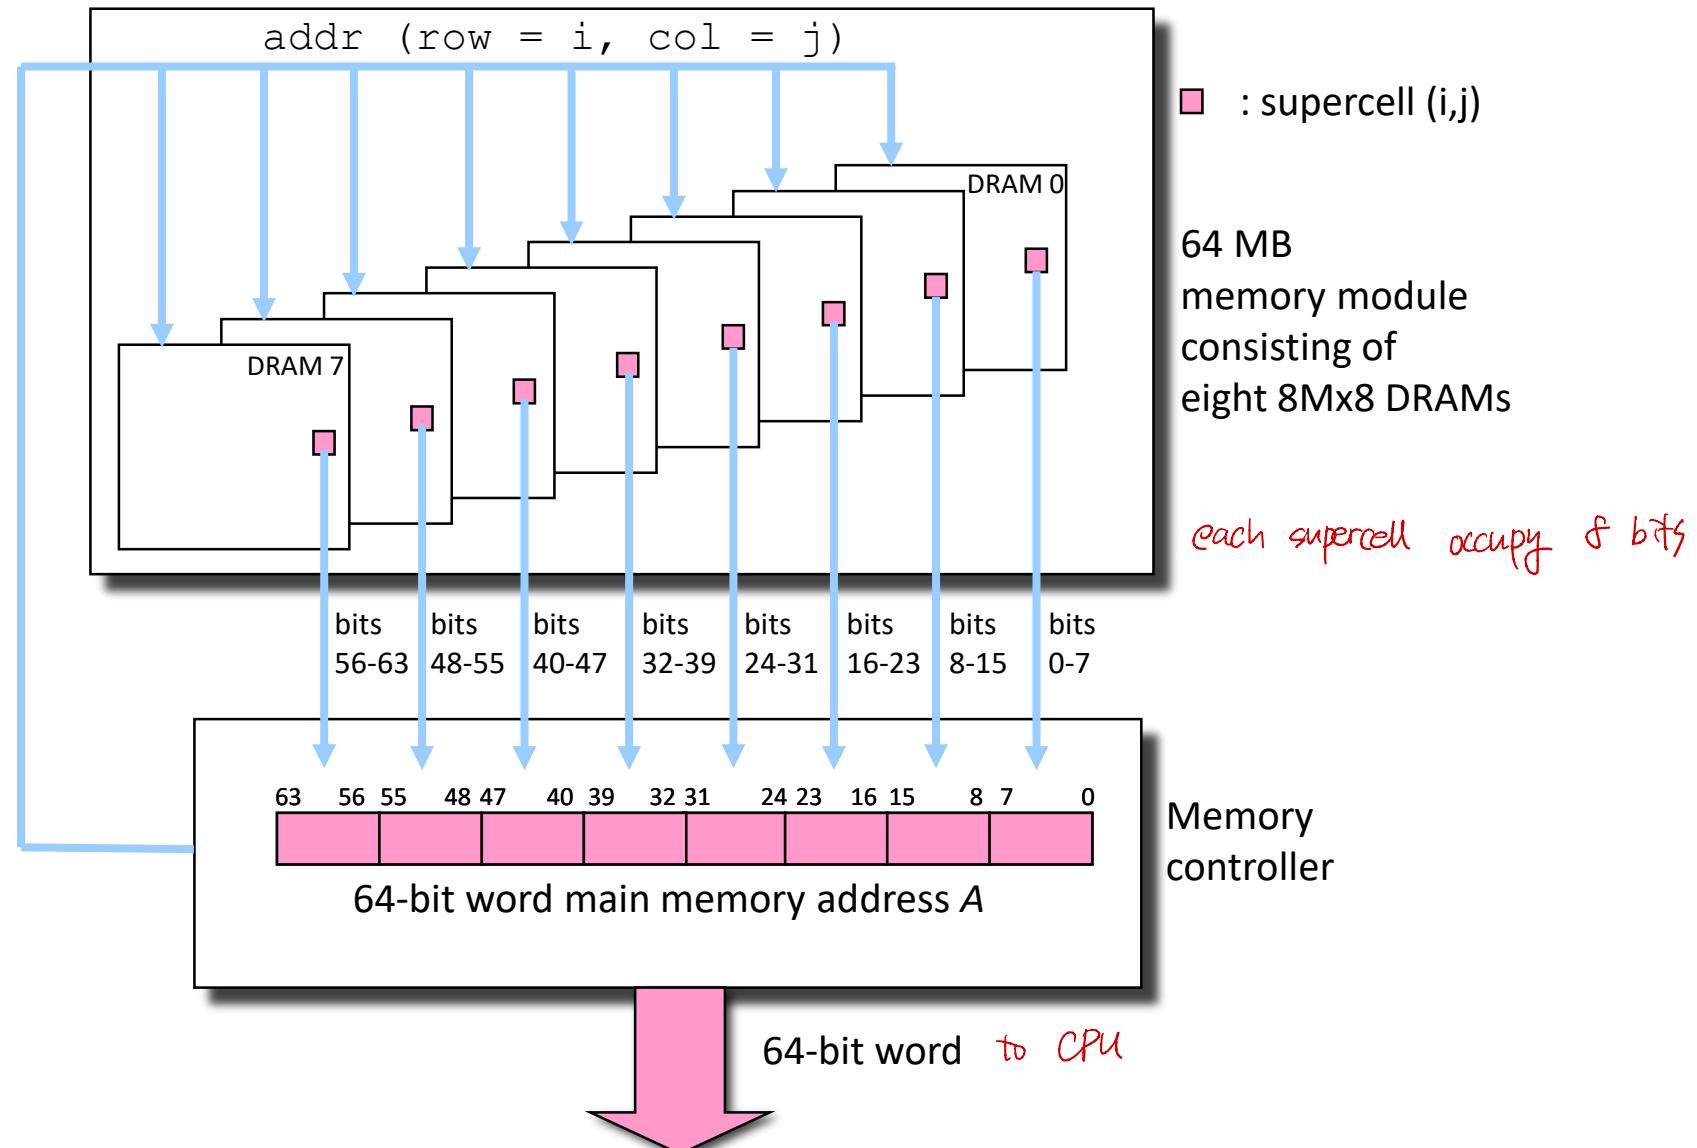
Step 2(a): Column access strobe (**CAS**) selects column 1.

Step 2(b): Supercell (2,1) copied from buffer to data lines, and eventually back to the CPU.



# Memory Modules

→ each module consists of multiple DRAM ⇒ access multiple data in parallel time



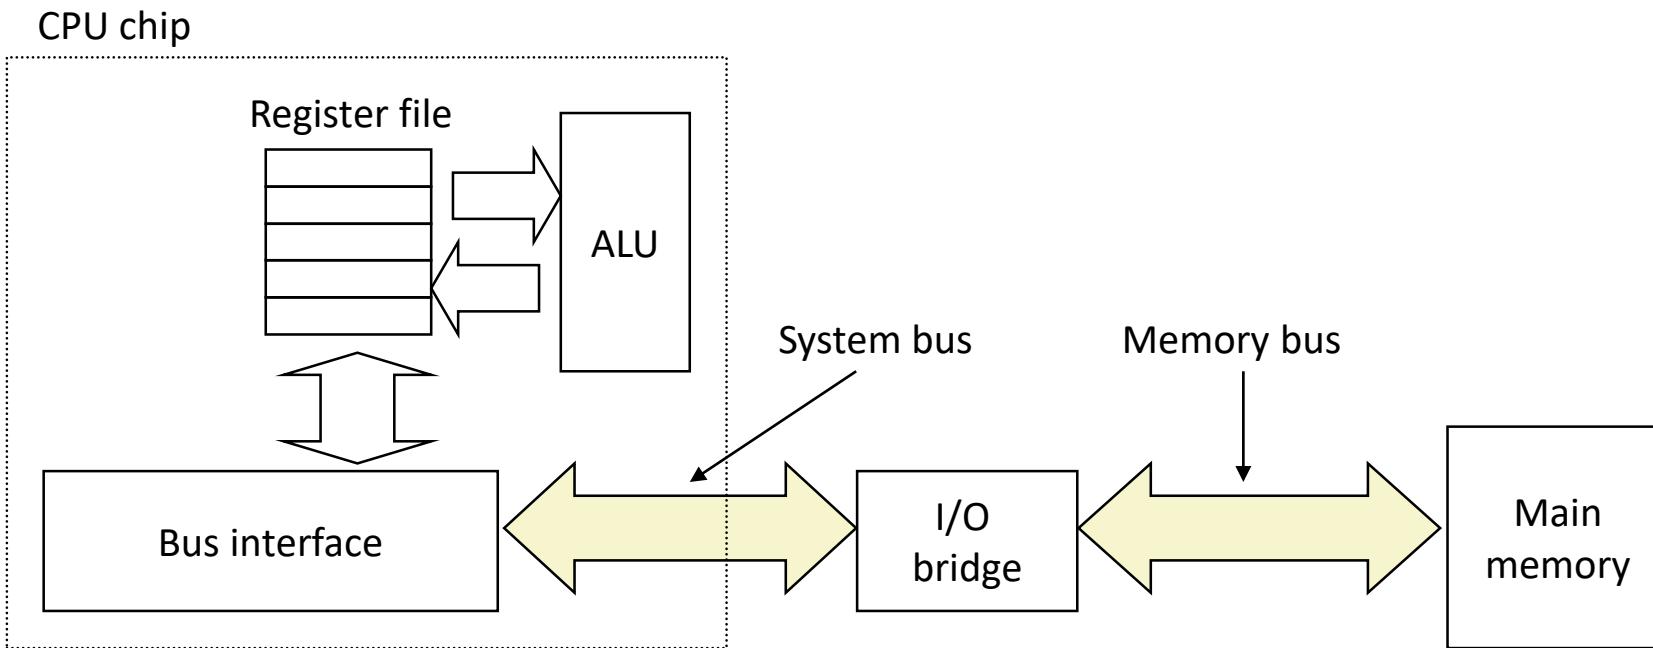
# Traditional Bus Structure Connecting CPU and Memory

- A **bus** is a collection of parallel wires that carry address, data, and control signals.

- Buses are typically shared by multiple devices.

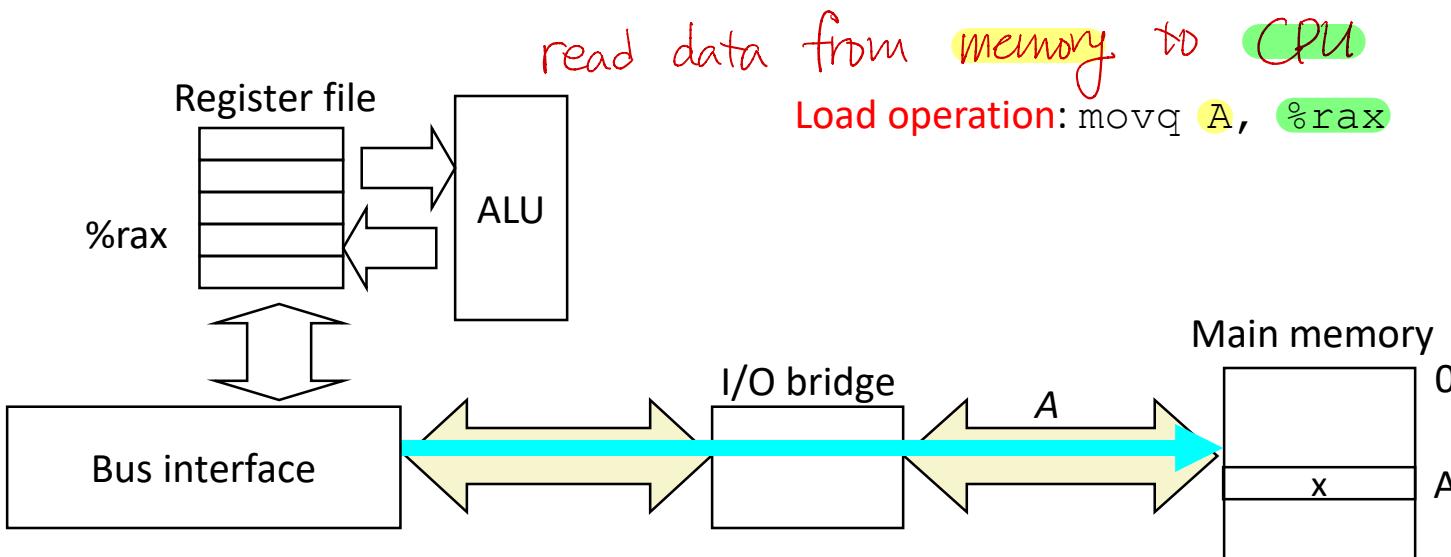
CPU ↔ Memory

such as I/O devices  
keyboard, mouse, ...



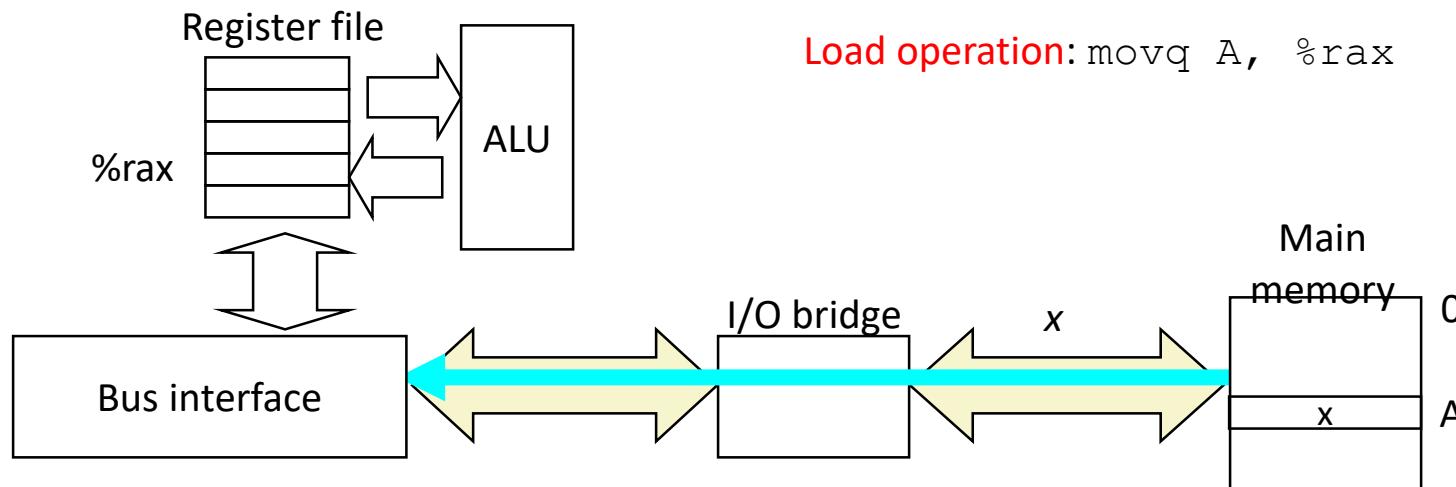
# Memory Read Transaction (1)

- CPU places address A on the memory bus.



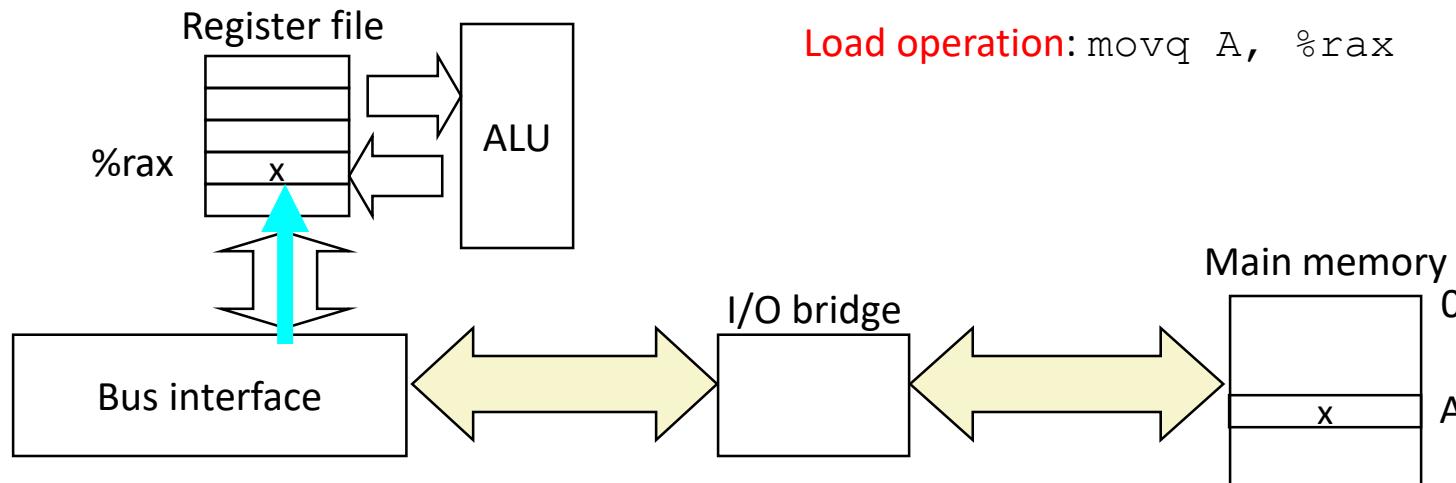
# Memory Read Transaction (2)

- Main memory reads A from the memory bus, retrieves word x, and places it on the bus.



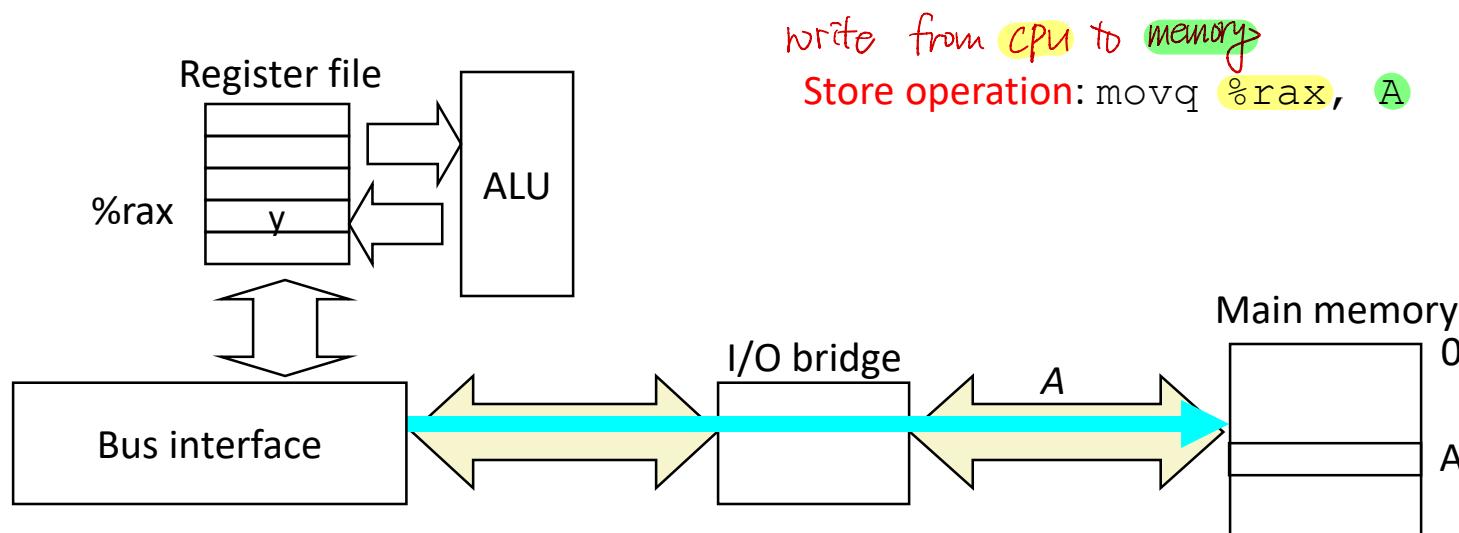
# Memory Read Transaction (3)

- CPU read word  $x$  from the bus and copies it into register  $\%rax$ .



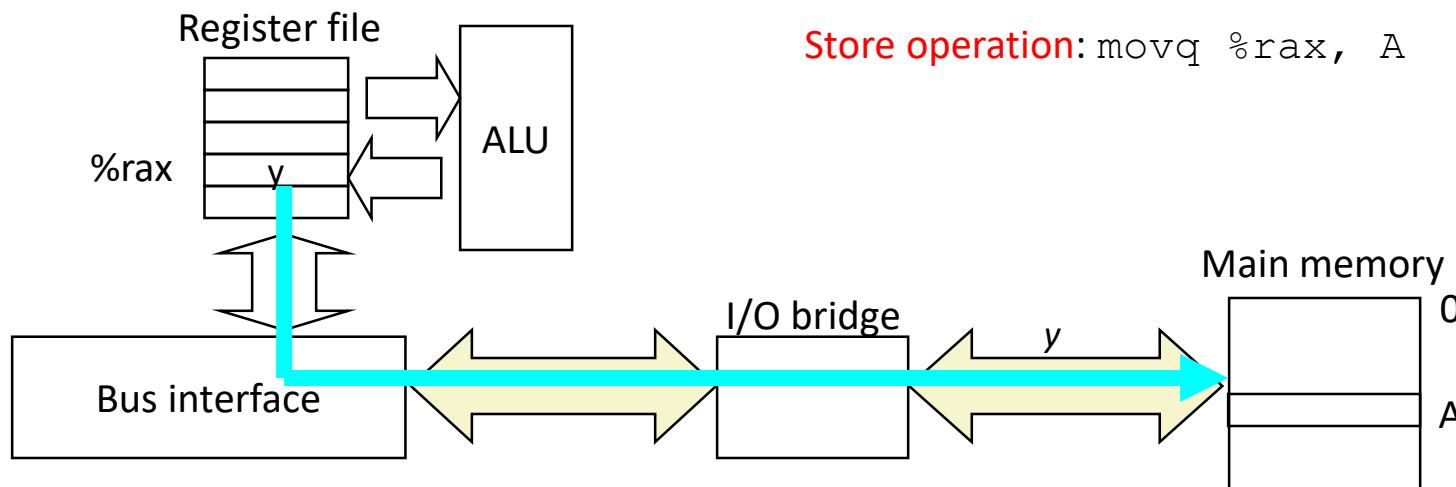
# Memory Write Transaction (1)

- CPU places address A on bus. Main memory reads it and waits for the corresponding data word to arrive.



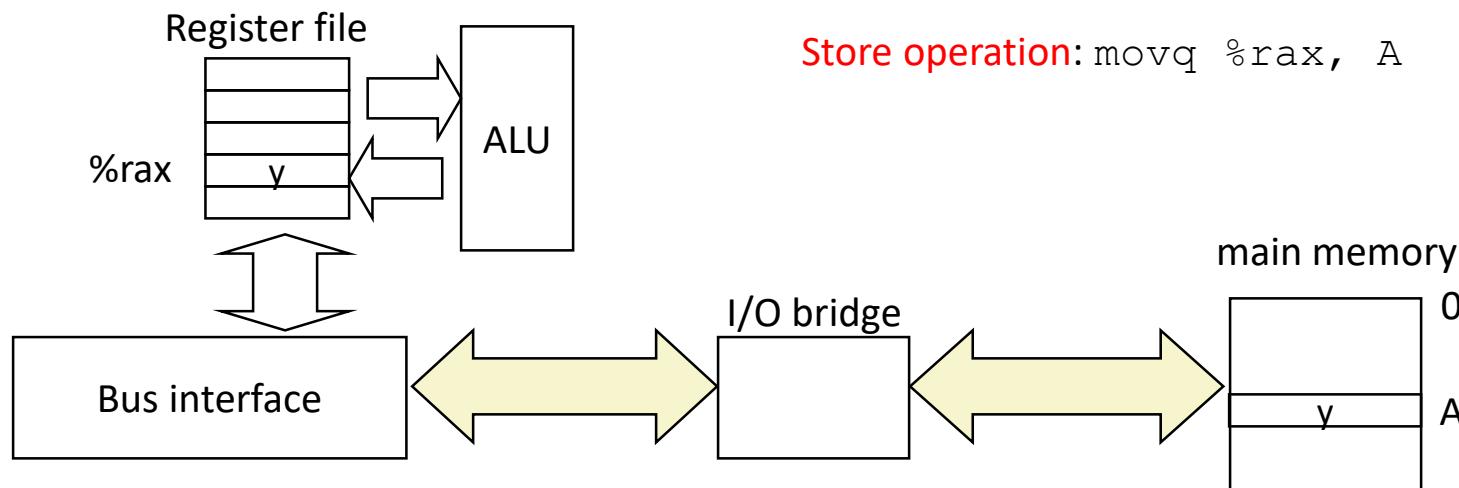
# Memory Write Transaction (2)

- CPU places data word  $y$  on the bus.



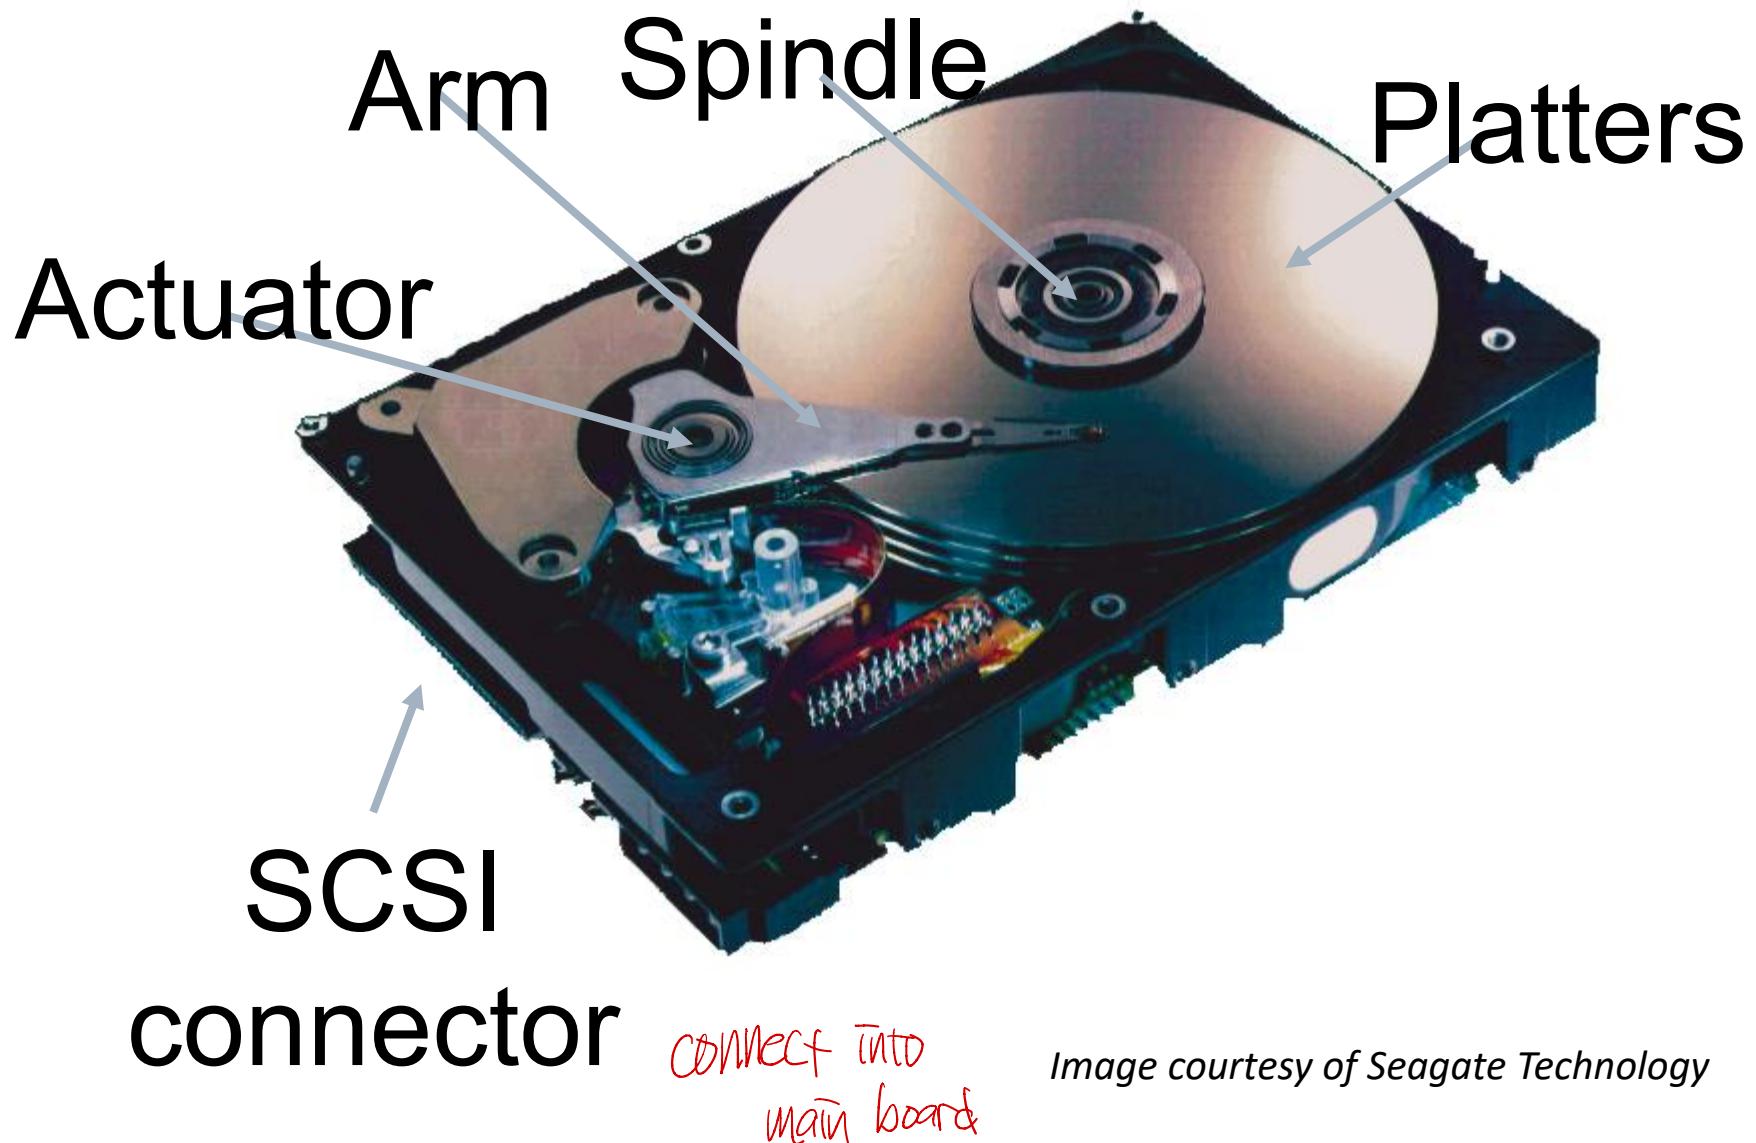
# Memory Write Transaction (3)

- Main memory reads data word  $y$  from the bus and stores it at address A.



# What's Inside A Disk Drive?

---

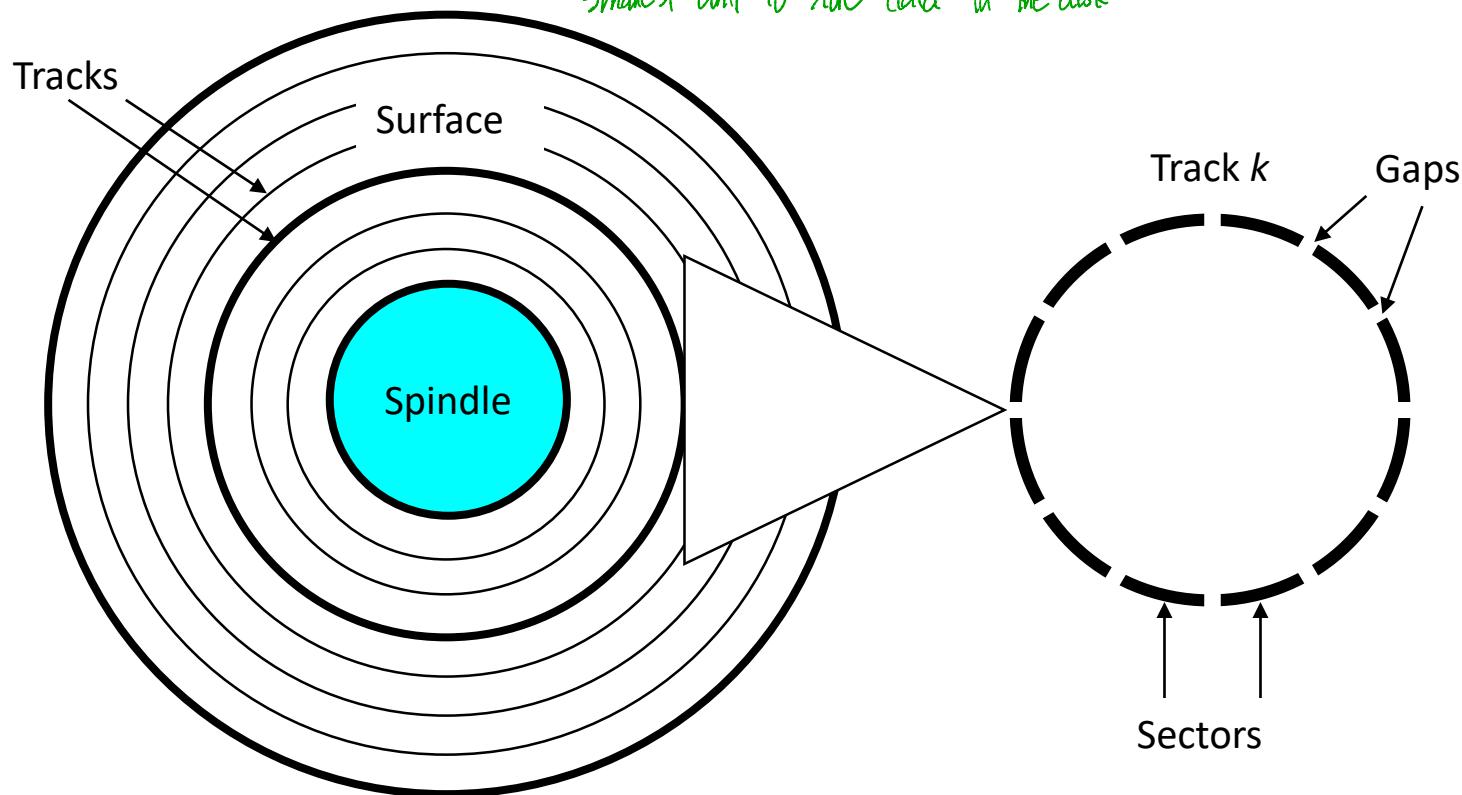


*Image courtesy of Seagate Technology*

# Disk Geometry

- Disks consist of **platters**, each with two **surfaces**. *front  
back*
- Each surface consists of concentric rings called **tracks**.
- Each track consists of **sectors** separated by **gaps**.

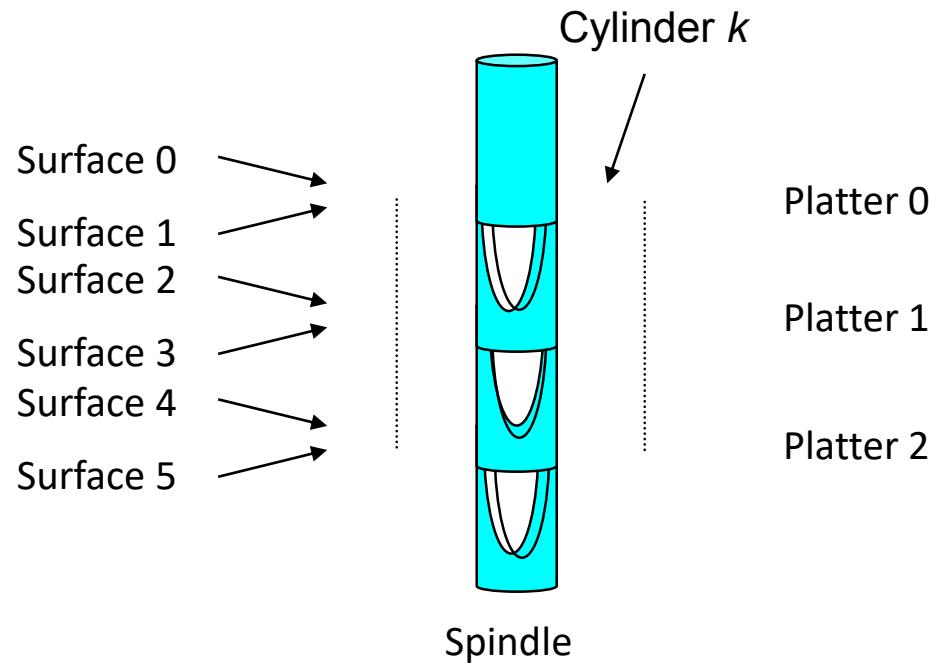
platter  
surface  
track  
sector



# Disk Geometry (Multiple-Platter View)

---

- Aligned tracks form a cylinder.



# Disk Capacity

---

## ■ **Capacity:** maximum number of bits that can be stored.

- Vendors express capacity in units of gigabytes (GB), where  
 $1 \text{ GB} = 10^9 \text{ Bytes}$ .

## ■ **Capacity is determined by these technology factors:**

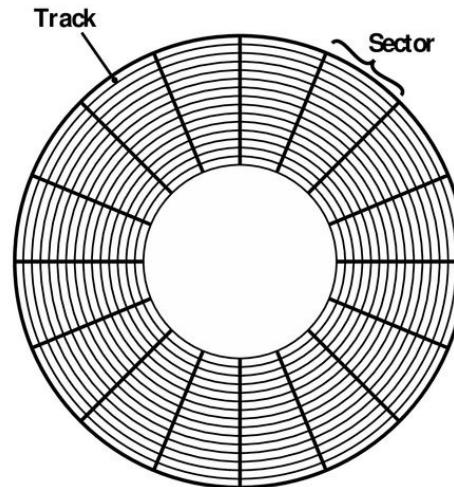
- **Recording density** (bits/in): number of bits that can be squeezed into a 1 inch segment of a track. *how much data in a track*
- **Track density** (tracks/in): number of tracks that can be squeezed into a 1 inch radial segment. *how many tracks in one surface*
- **Areal density** (bits/in<sup>2</sup>): product of recording and track density.

# Recording zones

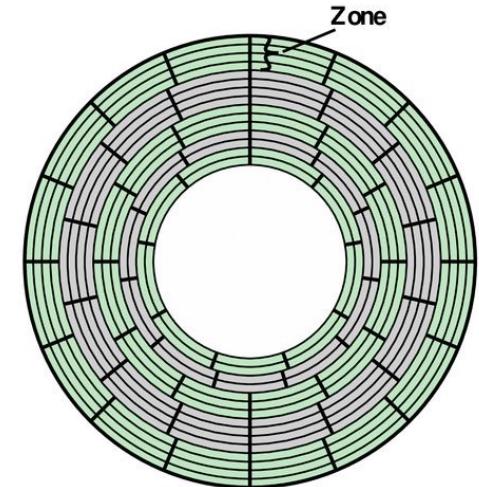
## ■ Modern disks partition tracks into disjoint subsets called **recording zones**

- Each track in a zone has the same number of sectors, determined by the circumference of innermost track.
- Each zone has a different number of sectors/track, outer zones have more sectors/track than inner zones.
- So we use **average** number of sectors /track when computing capacity.

**Figure 7.3**  
**Comparison of Disk Layout Methods**



(a) Constant angular velocity



(b) Multiple zone recording

# Computing Disk Capacity

Capacity = (# bytes/sector) x (avg. # sectors/track) x  
(# tracks/surface) x (# surfaces/platter) x  
(# platters/disk)

Example:

- 512 bytes/sector
- 300 sectors/track (on average)
- 20,000 tracks/surface
- 2 surfaces/platter
- 5 platters/disk

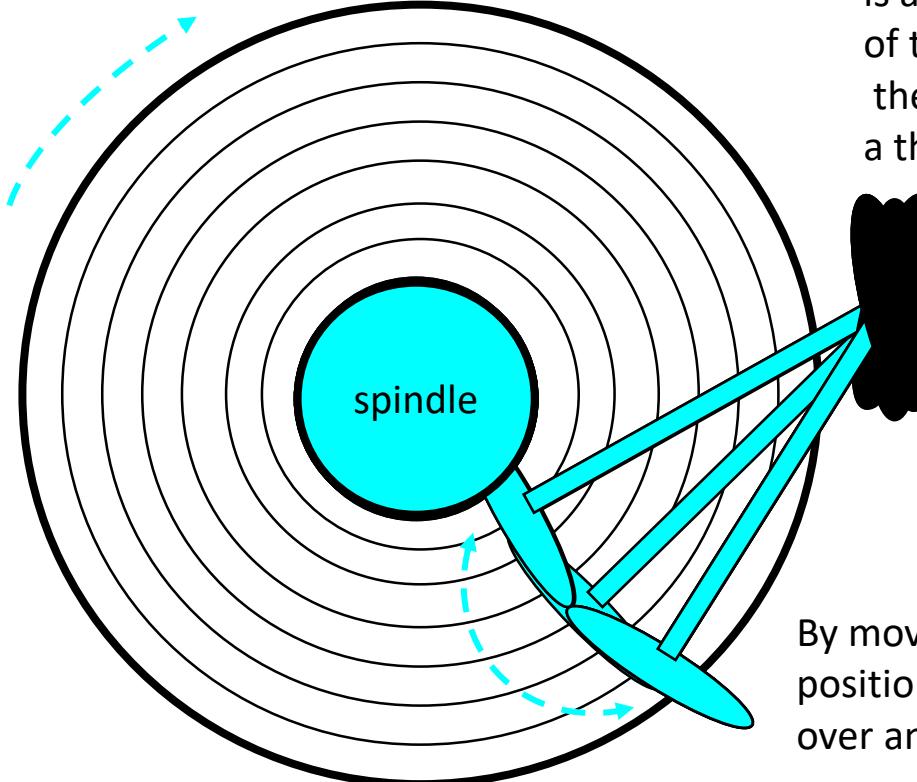
$$\Rightarrow \text{Capacity} = 512 \times 300 \times 20,000 \times 2 \times 5 = 30,720,000,000 \text{ bytes} = 30.72 \text{ GB}$$

cf)  $1\text{GB} = 10^9 \text{ byte}$

# Disk Operation (Single-Platter View)

---

The disk surface spins at a fixed rotational rate

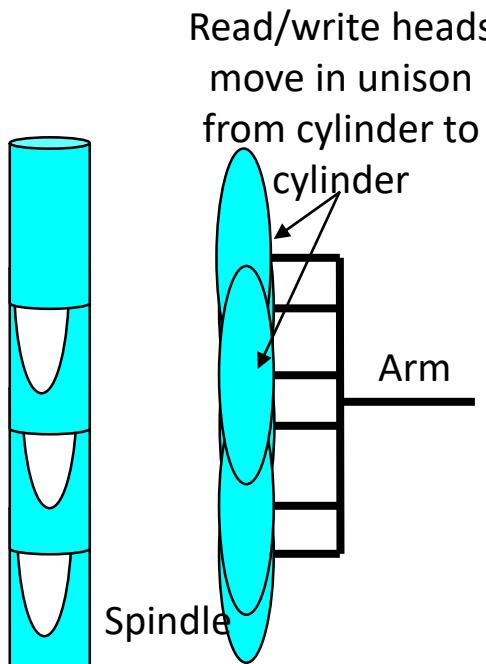
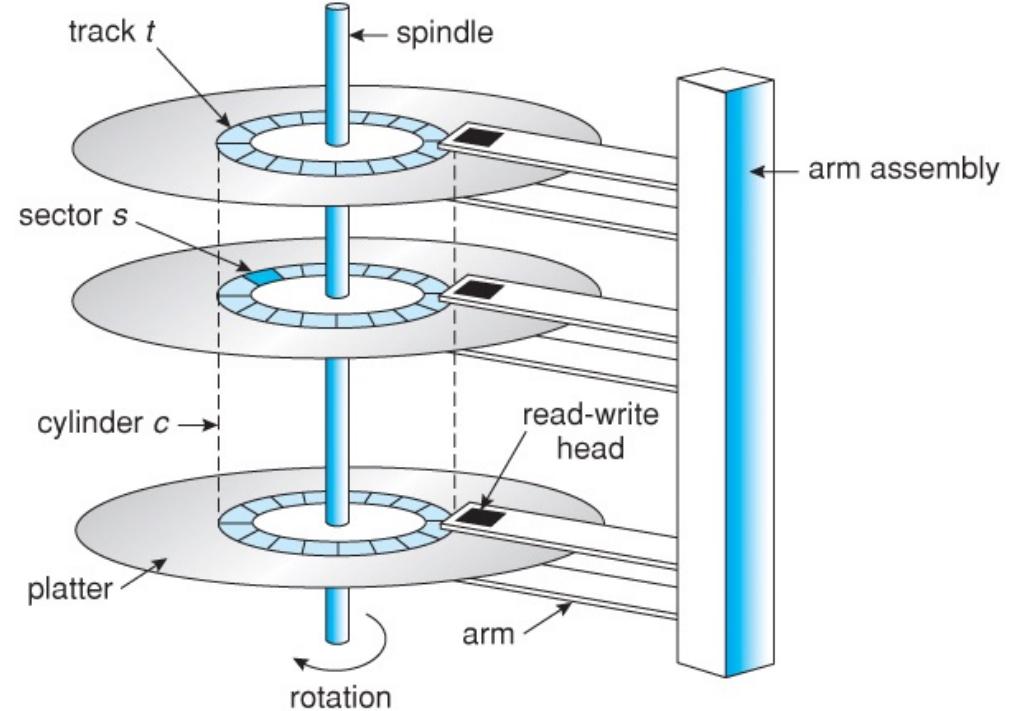


The read/write *head* is attached to the end of the *arm* and flies over the disk surface on a thin cushion of air.

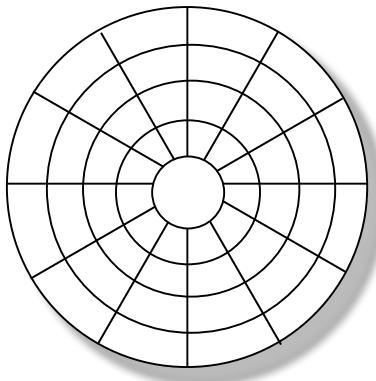
By moving radially, the arm can position the read/write head over any track.

# Disk Operation (Multi-Platter View)

---



# Disk Structure - top view of single platter

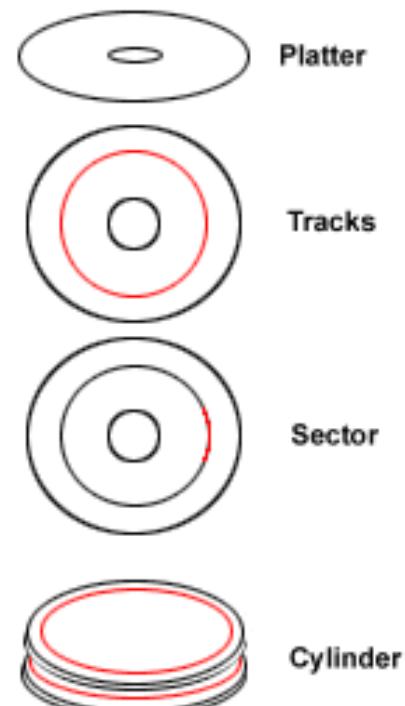
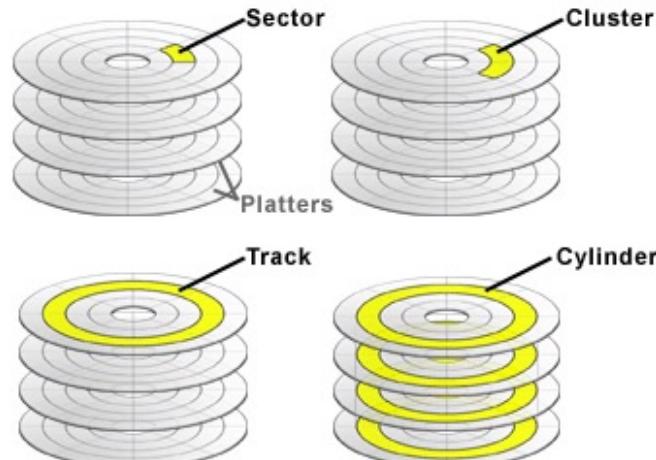


ARM can access multiple tracks  
in the same cylinder at a time!

Surface organized into tracks

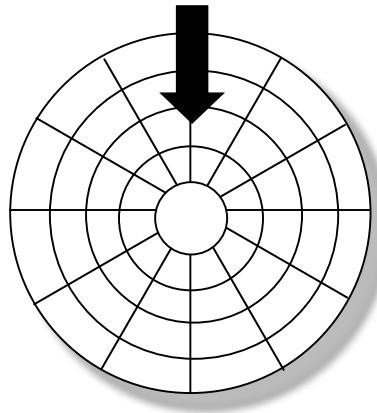
Tracks divided into sectors

Cylinder  $\Rightarrow$  set of multiple tracks for  
all the platters .



# Disk Access

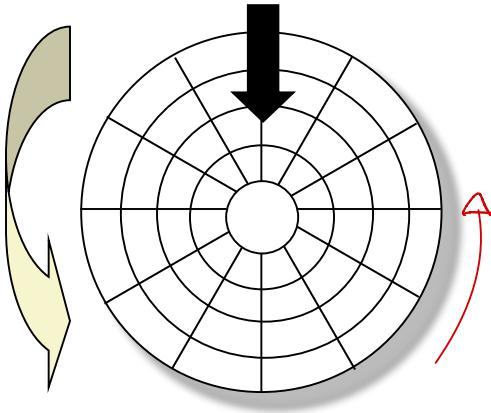
---



Head in position above a track

# Disk Access

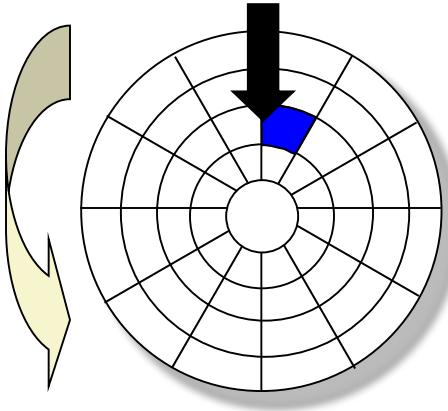
---



Rotation is counter-clockwise  
*always!*

# Disk Access – Read

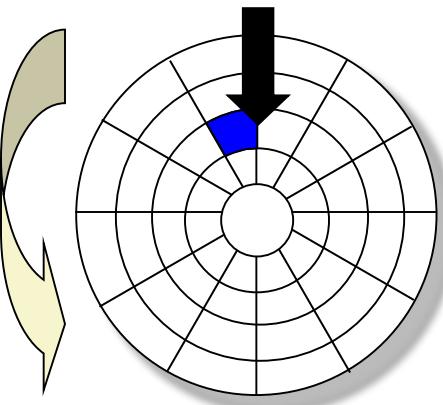
---



About to read blue sector

# Disk Access – Read

---

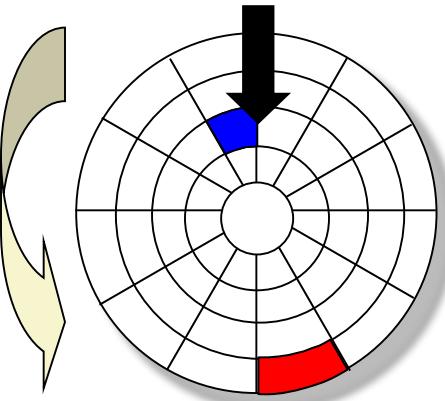


After BLUE read

After reading blue sector

# Disk Access – Read

---

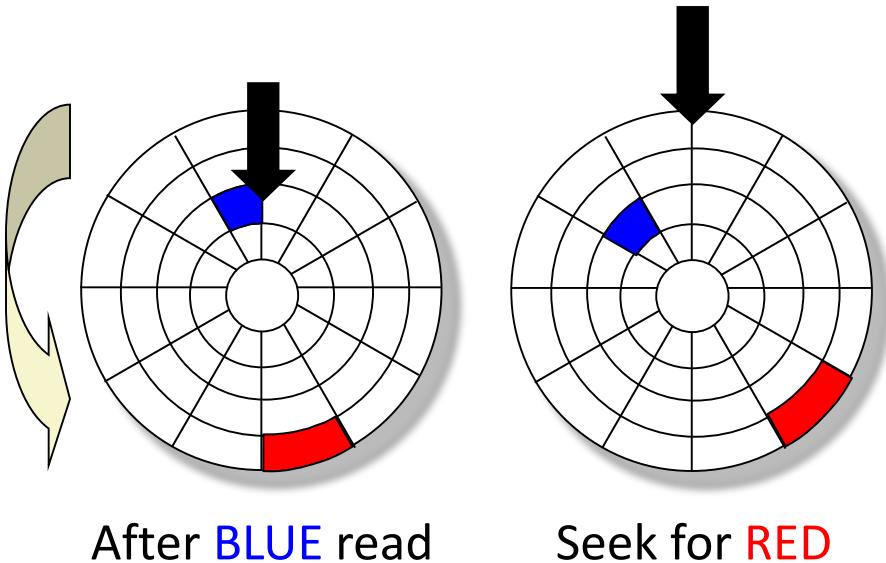


After **BLUE** read

Red request scheduled next

# Disk Access – Seek

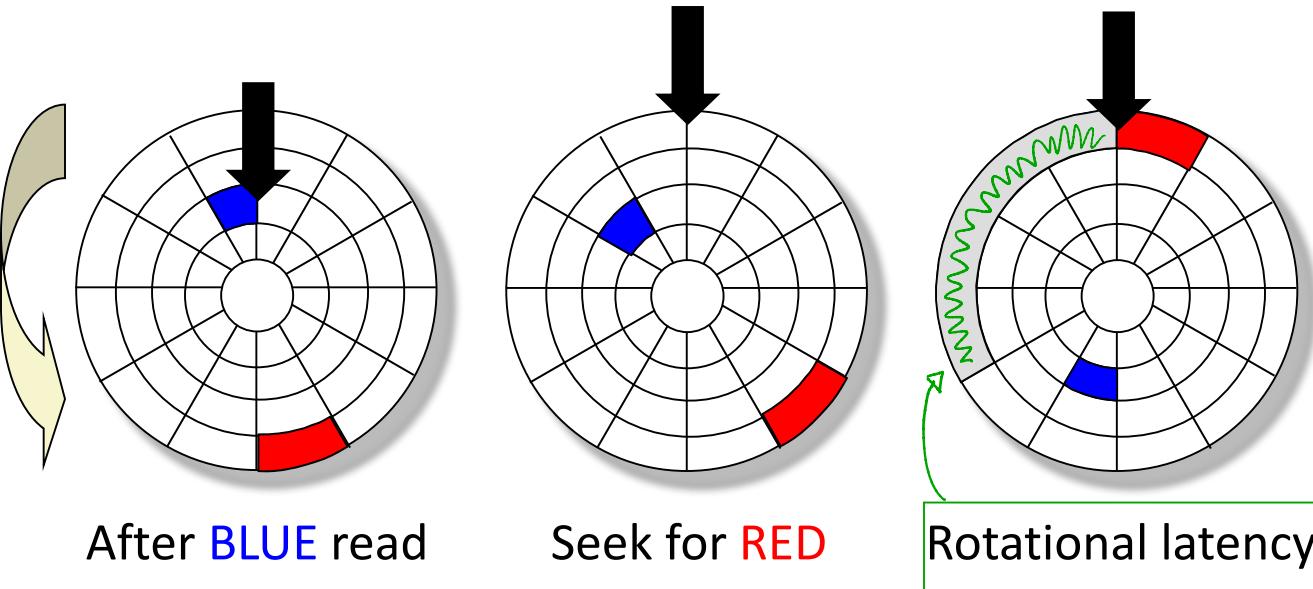
---



Seek to red's track

# Disk Access – Rotational Latency

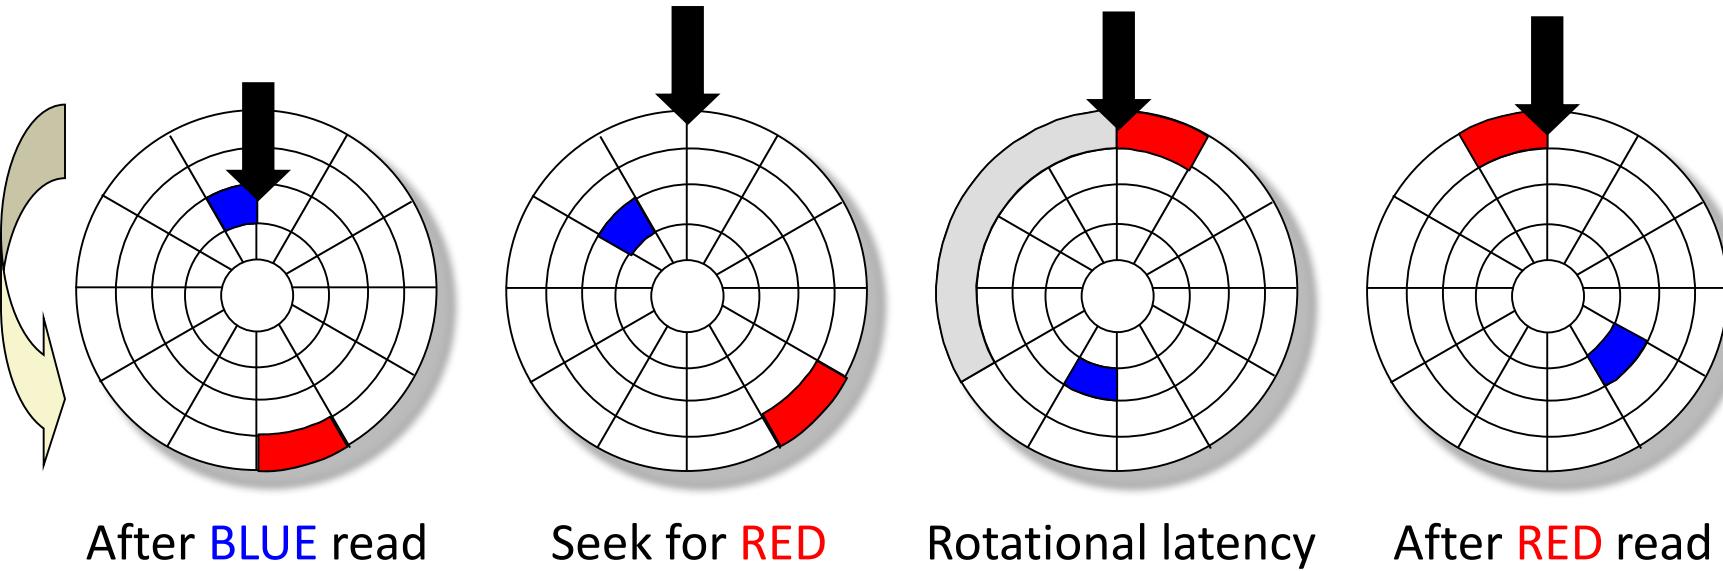
---



Wait for red sector to rotate around

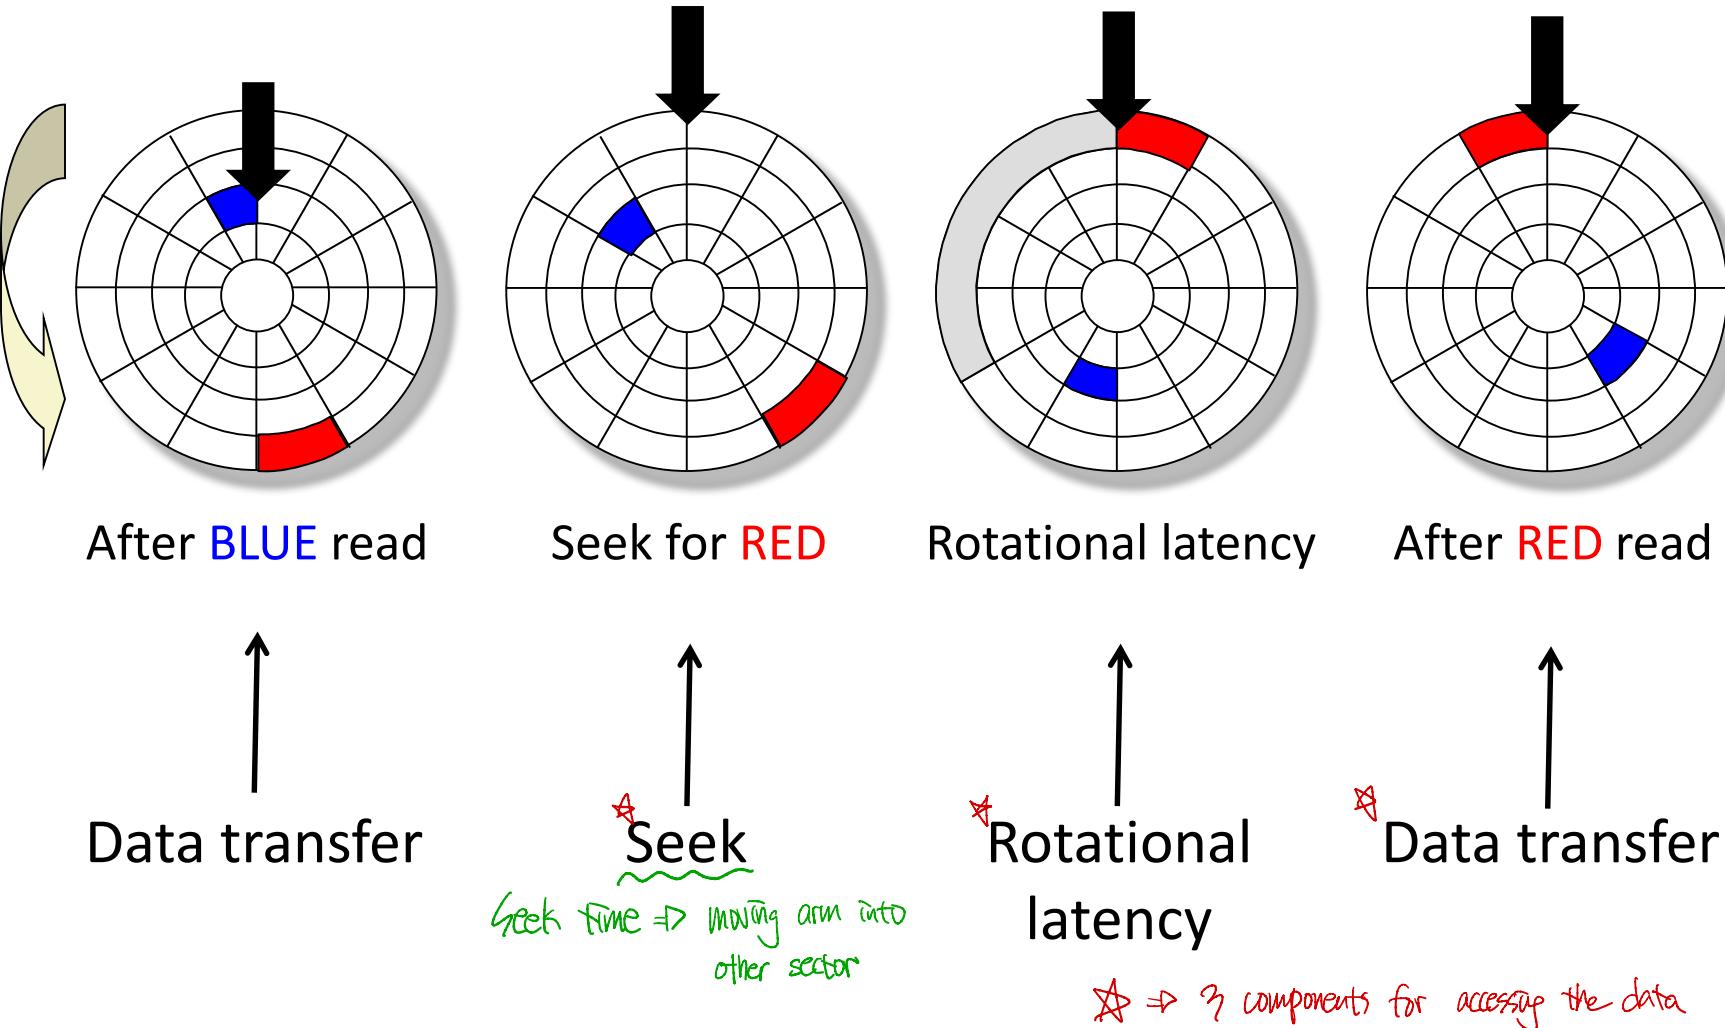
# Disk Access – Read

---



Complete read of red

# Disk Access – Service Time Components



# Disk Access Time

---

## ■ Average time to access some target sector approximated by :

- $T_{\text{access}} = T_{\text{avg seek}} + T_{\text{avg rotation}} + T_{\text{avg transfer}}$

## ■ Seek time ( $T_{\text{avg seek}}$ )

- Time to position heads over cylinder containing target sector.
- Typical  $T_{\text{avg seek}}$  is 3—9 ms

## ■ Rotational latency ( $T_{\text{avg rotation}}$ )

- Time waiting for first bit of target sector to pass under r/w head.
- $T_{\text{avg rotation}} = 1/2 \times 1/\text{RPMs} \times 60 \text{ sec}/1 \text{ min}$
- Typical  $T_{\text{avg rotation}} = 7200 \text{ RPMs}$

## ■ Transfer time ( $T_{\text{avg transfer}}$ )

- Time to read the bits in the target sector.
- $T_{\text{avg transfer}} = 1/\text{RPM} \times 1/(\text{avg # sectors/track}) \times 60 \text{ secs}/1 \text{ min.}$

# Disk Access Time Example

## Given:

- Rotational rate = 7,200 RPM
- Average seek time = 9 ms.
- Avg # sectors/track = 400.

$$1\text{ sec} = 1000 \text{ ms}$$

## Derived:

- $T_{\text{avg rotation}} = \left(\frac{1}{2} \times \frac{1}{7200} \times 60\right) \times 1000 \text{ ms/sec} = 4 \text{ ms}$
- $T_{\text{avg transfer}} = \left(\frac{1}{7200} \times \frac{1}{400} \times 60\right) \times 1000 \text{ ms/sec} \approx 0.02 \text{ ms}$
- $T_{\text{access}} = 9 \text{ ms} + 4 \text{ ms} + 0.02 \text{ ms} = 13.02 \text{ ms}$

## Important points:

- Access time dominated by seek time and rotational latency.
- First bit in a sector is the most expensive, the rest are free.
- SRAM access time is about 4 ns/doubleword, DRAM about 60 ns
  - Disk is about 40,000 times slower than SRAM,
  - 2,500 times slower than DRAM.