

Computation

Sim, Min Kyu, Ph.D., mksim@seoultech.ac.kr



The first half of our course

The second half of our course

About

- This lecture note covers how to use statistical software, R, to deal with computations in linear algebra.
- Installation and basic guides for R can be found at:
https://github.com/aceMKSIm/teaching/blob/master/Data%20Visualization/Lecture%20Notes/L01_installR.pdf
- More interested in R and data visualization?
 - You can access my lecture note for a graduate course, 'Data Visualization' at <https://github.com/aceMKSIm/teaching/tree/master/Data%20Visualization>
 - You can request an invitation link to a superb online course platform, www.datacamp.com, where courses for R and Python are available free through my invitation.

The first half of our course

Matrix Input

- Matrix input is done by 1) listing all elements by `c()` function, 2) telling the exact number of row by `nrow=`, and 3) whether you want to list elements by row or by column (`byrow`).

```
A1 <- matrix(c(1,-3,0,2,-1,5,1,2,3),
              nrow = 3, byrow = FALSE)
b1 <- matrix(c(0,1,-1),
              nrow=3, byrow=FALSE)
A2 <- matrix(c(4,2,2,-1,1,-1,6,6,8),
              nrow=3, byrow=FALSE)
A3 <- matrix(c(1,0,-1,-7,0,7,0,1,4,6,-2,2),
              nrow=3, byrow=FALSE)
b3 <- matrix(c(5,-3,7),
              nrow=3, byrow=FALSE)
# Following to check for correct input
A1
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]   -3   -1    2
## [3,]    0    5    3
```

b1

```
##      [,1]
## [1,]    0
## [2,]    1
## [3,]   -1
```

A2

```
##      [,1] [,2] [,3]
## [1,]    4   -1    6
## [2,]    2    1    6
## [3,]    2   -1    8
```

A3

```
##      [,1] [,2] [,3] [,4]
## [1,]    1   -7    0    6
## [2,]    0    0    1   -2
## [3,]   -1    7    4    2
```

b3

```
##      [,1]
## [1,]    5
## [2,]   -3
## [3,]    7
```

- Concatenating function `cbind()` allows to construct a augmented matrix.

```
Aug3 <- cbind(A3, b3)
```

Aug3

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1  -7    0    6    5
## [2,]    0    0    1   -2   -3
## [3,]   -1    7    4    2    7
```

Reduced Echelon Form

- Package (often called as library) provides additional functionality.
- In the first ever to use a package, it must be installed by following syntax.

```
install.packages("pracma")
```

- In your computation session, you must use following syntax to activate the package.

```
library(pracma)
```

- The package pracma offers rref() that gives the reduced echelon form.

```
library(pracma)
```

```
A1
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]   -3   -1    2
## [3,]    0    5    3
```

```
rref(A1)
```

```
##      [,1] [,2] [,3]
## [1,]    1    0    0
## [2,]    0    1    0
## [3,]    0    0    1
```

A2

```
##      [,1] [,2] [,3]
## [1,]    4  -1    6
## [2,]    2   1    6
## [3,]    2  -1    8
```

rref(A2)

```
##      [,1] [,2] [,3]
## [1,]    1   0   0
## [2,]    0   1   0
## [3,]    0   0   1
```

A3

```
##      [,1] [,2] [,3] [,4]
## [1,]    1  -7   0   6
## [2,]    0   0   1  -2
## [3,]   -1   7   4   2
```

rref(A3)

```
##      [,1] [,2] [,3] [,4]
## [1,]    1  -7   0   0
## [2,]    0   0   1   0
## [3,]    0   0   0   1
```

Aug3

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1  -7   0   6   5
## [2,]    0   0   1  -2  -3
## [3,]   -1   7   4   2   7
```

rref(Aug3)

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1  -7   0   0 -4.0
## [2,]    0   0   1   0  0.0
## [3,]    0   0   0   1  1.5
```


Solving $Ax = b$

```
A1
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]   -3   -1    2
## [3,]    0    5    3
```

```
b1
##      [,1]
## [1,]    0
## [2,]    1
## [3,]   -1
```

```
x <- solve(A1, b1)
```

```
x
##      [,1]
## [1,]  0.6
## [2,] -0.8
## [3,]  1.0
```

- `%%` is an operator for matrix multiplication.

```
# Sanity check
```

```
A1 %% x
##      [,1]
## [1,]    0
## [2,]    1
## [3,]   -1
b1
```

```
##      [,1]
## [1,]    0
## [2,]    1
## [3,]   -1
```

Inverting

A1

```
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]   -3   -1    2
## [3,]    0    5    3
```

```
inv_A1 <- solve(A1)
```

inv_A1

```
##      [,1] [,2] [,3]
## [1,]  1.3  0.1 -0.5
## [2,] -0.9 -0.3  0.5
## [3,]  1.5  0.5 -0.5
```

Sanity check

A1 %*% inv_A1

```
##      [,1] [,2] [,3]
## [1,]    1    0    0
## [2,]    0    1    0
## [3,]    0    0    1
```

inv_A1 %*% A1

```
##      [,1] [,2]      [,3]
## [1,]    1    0 0.000000e+00
## [2,]    0    1 2.220446e-16
## [3,]    0    0 1.000000e+00
```

Rank

```
library(Matrix)
```

```
A1
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]   -3   -1    2
## [3,]    0    5    3
```

```
rankMatrix(A1)
```

```
## [1] 3
## attr(,"method")
## [1] "tolNorm2"
## attr(,"useGrad")
## [1] FALSE
## attr(,"tol")
## [1] 6.661338e-16
```

LU decomposition

```
library(matrixcalc)
```

```
A1
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]   -3   -1    2
## [3,]    0    5    3
```

```
lu.decomposition(A1)
```

```
## $L
```

```
##      [,1] [,2] [,3]
## [1,]    1    0    0
## [2,]   -3    1    0
## [3,]    0    1    1
```

```
##
```

```
## $U
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]    0    5    5
## [3,]    0    0   -2
```

```
# Sanity check
```

```
A1_L <- lu.decomposition(A1)$L
```

```
A1_U <- lu.decomposition(A1)$U
```

```
A1_L %*% A1_U
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]   -3   -1    2
## [3,]    0    5    3
```

```
A1
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]   -3   -1    2
## [3,]    0    5    3
```

Determinants

A1

```
##      [,1] [,2] [,3]
## [1,]    1    2    1
## [2,]   -3   -1    2
## [3,]    0    5    3
```

`det(A1)`

```
## [1] -10
```

A2

```
##      [,1] [,2] [,3]
## [1,]    4   -1    6
## [2,]    2    1    6
## [3,]    2   -1    8
```

`det(A2)`

```
## [1] 36
```

The second half of our course

Eigenvalue and Eigenvector

A2

```
##      [,1] [,2] [,3]
## [1,]    4  -1    6
## [2,]    2   1    6
## [3,]    2  -1    8
```

`eigen(A2)`

```
## eigen() decomposition
## $values
## [1] 9 2 2
##
## $vectors
##      [,1]      [,2]      [,3]
## [1,] 0.5773503 0.2915782 0.3377822
## [2,] 0.5773503 0.8863713 0.0000000
## [3,] 0.5773503 0.0505358 -0.1125941
```

`eigen(A2)$values`

```
## [1] 9 2 2
```

`eigen(A2)$vectors`

```
##      [,1]      [,2]      [,3]
## [1,] 0.5773503 0.2915782 0.3377822
## [2,] 0.5773503 0.8863713 0.0000000
## [3,] 0.5773503 0.0505358 -0.1125941
```

```
lambda1 <- eigen(A2)$values[1]
lambda2 <- eigen(A2)$values[2]
lambda3 <- eigen(A2)$values[3]
v1 <- eigen(A2)$vectors[,1]
v2 <- eigen(A2)$vectors[,2]
v3 <- eigen(A2)$vectors[,3]
```

```
lambda1
## [1] 9
lambda2
## [1] 2
lambda3
## [1] 2
v1
## [1] 0.5773503 0.5773503 0.5773503
v2
## [1] 0.2915782 0.8863713 0.0505358
v3
## [1] 0.3377822 0.0000000 -0.1125941
```


Diagonalization

$$A = PDP^{-1}$$

```
P <- eigen(A2)$vectors
D <- diag(eigen(A2)$values)
```

```
P
##          [,1]      [,2]      [,3]
## [1,] 0.5773503 0.2915782 0.3377822
## [2,] 0.5773503 0.8863713 0.0000000
## [3,] 0.5773503 0.0505358 -0.1125941
```

```
D
##          [,1] [,2] [,3]
## [1,]      9    0    0
## [2,]      0    2    0
## [3,]      0    0    2
```

```
# sanity check
```

```
A2
##          [,1] [,2] [,3]
## [1,]      4   -1    6
## [2,]      2    1    6
## [3,]      2   -1    8
```

```
P %*% D %*% solve(P)
```

```
##          [,1] [,2] [,3]
## [1,]      4   -1    6
## [2,]      2    1    6
## [3,]      2   -1    8
```

Dataset *mtcars*

- The rest of this lecture note will use the one of the popular built-in datasets in R, *mtcars*, which includes several cars and their fuel efficiency. To analyze the dataset.
- Typing `help(mtcars)` in your console will provide basic description on the dataset.

```
help(mtcars)
```

The data was extracted from the 1974 Motor Trend US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973–74 models)

- The dataset can be explored by following syntax, `str()`.

```
str(mtcars)
```

```
## 'data.frame':   32 obs. of  11 variables:
## $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
## $ cyl : num   6  6  4  6  8  6  8  4  4  6 ...
## $ disp: num  160 160 108 258 360 ...
## $ hp  : num  110 110 93 110 175 105 245 62 95 123 ...
## $ drat: num   3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
## $ wt  : num   2.62 2.88 2.32 3.21 3.44 ...
## $ qsec: num  16.5 17 18.6 19.4 17 ...
## $ vs  : num   0  0  1  1  0  1  0  1  1  1 ...
## $ am  : num   1  1  1  0  0  0  0  0  0  0 ...
## $ gear: num   4  4  4  3  3  3  3  4  4  4 ...
## $ carb: num   4  4  1  1  2  1  4  2  2  4 ...
```

Each variable's description is also provided from the help page as follows.

- [1] mpg, Miles/(US) gallon
- [2] cyl, Number of cylinders
- [3] disp, Displacement (cu.in.)
- [4] hp, Gross horsepower
- [5] drat, Rear axle ratio
- [6] wt, Weight (1000 lbs)
- [7] qsec, 1/4 mile time
- [8] vs, Engine (0 = V-shaped, 1 = straight)
- [9] am, Transmission (0 = automatic, 1 = manual)
- [10] gear, Number of forward gears
- [11] carb, Number of carburetors
- The first variable mpg refers to the fuel efficiency, since it measures how many miles the car can move with a gallon of its fuel. The other 10 variables should affect the mpg. That is, we call mpg as a dependent variable and the other variables as independent variables.

Linear regression

- For simplicity, the rest of this will use only the first six independent variables from now on: cyl, disp, hp, drat, wt, qsec
- We will use
 - disp as an independent variable (i.e. 'X' variable)
 - mpg as dependent variable (i.e. 'Y' variable).
 - Following codes will reveal a few observed pairs of (X, Y) .
 - i.e. the function head() reveals the first six observations of the dataset.

```
head(mtcars[,c("disp", "mpg")])
```

```
##           disp  mpg
## Mazda RX4    160 21.0
## Mazda RX4 Wag 160 21.0
## Datsun 710    108 22.8
## Hornet 4 Drive 258 21.4
## Hornet Sportabout 360 18.7
## Valiant      225 18.1
```

- Based on the textbook (Ch.6, Example 4),
 - we shall construct a A matrix where
 - the first column is one-vector (generated by `rep(1, length(mtcars$disp))`)
 - the second column is the observed quantity of `disp` (generated by `mtcars$disp`).
 - The matrix (or, a vector) b is nothing but the quantities of `mtcars$mpg`.
- This can be done with following codes. (The function `matrix()` is another way to create matrix.)

```
A <- matrix(cbind(rep(1,length(mtcars$disp)), mtcars$disp), ncol = 2)
b <- matrix(mtcars$mpg, ncol = 1)
```

- To make sure A and b are correctly assigned, following codes reveal its beginning part.

`head(A)`

```
##      [,1] [,2]
## [1,]    1 160
## [2,]    1 160
## [3,]    1 108
## [4,]    1 258
## [5,]    1 360
## [6,]    1 225
```

`head(b)`

```
##      [,1]
## [1,] 21.0
## [2,] 21.0
## [3,] 22.8
## [4,] 21.4
## [5,] 18.7
## [6,] 18.1
```

- Using A and b , the least-squares solution is obtained by solving

$$A^t Ax = A^t b$$

- Assign $(A^t A)^{-1} A^t b$ to the variable x

```
x <- solve(t(A) %*% A) %*% t(A) %*% b
```

```
x
```

```
##           [,1]
```

```
## [1,] 29.59985476
```

```
## [2,] -0.04121512
```

```
x[1] # alpha
```

```
## [1] 29.59985
```

```
x[2] # beta
```

```
## [1] -0.04121512
```


- As a statistical software, R offers linear regression function called `lm()`.
- Basic usage is given below. Confirm that the results match the numbers in the previous page.

```
lm(mtcars$mpg ~ mtcars$disp)
```

```
##  
## Call:  
## lm(formula = mtcars$mpg ~ mtcars$disp)  
##  
## Coefficients:  
## (Intercept)  mtcars$disp  
##      29.59985      -0.04122
```


- Following is visualization of original data and regression line.

```
library(ggplot2)
ggplot(mtcars, aes(x=displacement, y=mpg)) +
  geom_point() +
  geom_abline(intercept = x[1], slope = x[2])
```



Principal Component Analysis

Preparation

- PCA (Principal Component Analysis) is to identify common underlying components in dataset.
- In our set of six independent variables, i.e. `mtcars[,c("cyl", "disp", "hp", "drat", "wt", "qsec")]`, the six variables are naturally highly correlated to each other.
- For example, larger engine size (`disp`) is likely to be positively correlated to its horse power (`hp`) and its weight (`wt`).
- Let X be the dataset of the independent variables.

```
X <- as.matrix(mtcars[,c("cyl", "disp", "hp", "drat", "wt", "qsec")])  
head(X)
```

##		cyl	disp	hp	drat	wt	qsec
##	Mazda RX4	6	160	110	3.90	2.620	16.46
##	Mazda RX4 Wag	6	160	110	3.90	2.875	17.02
##	Datsun 710	4	108	93	3.85	2.320	18.61
##	Hornet 4 Drive	6	258	110	3.08	3.215	19.44
##	Hornet Sportabout	8	360	175	3.15	3.440	17.02
##	Valiant	6	225	105	2.76	3.460	20.22

Correlation matrix

- To explore its correlation structure, the following code generates its correlation matrix.

```
cor(X)
```

```
##          cyl      disp      hp      drat      wt      qsec
## cyl    1.0000000  0.9020329  0.8324475 -0.69993811  0.7824958 -0.59124207
## disp  0.9020329  1.0000000  0.7909486 -0.71021393  0.8879799 -0.43369788
## hp    0.8324475  0.7909486  1.0000000 -0.44875912  0.6587479 -0.70822339
## drat -0.6999381 -0.7102139 -0.4487591  1.00000000 -0.7124406  0.09120476
## wt    0.7824958  0.8879799  0.6587479 -0.71244065  1.0000000 -0.17471588
## qsec -0.5912421 -0.4336979 -0.7082234  0.09120476 -0.1747159  1.00000000
```

Assumption of PCA

- PCA assumes that, each observation can be tersely expressed without using all of the six variables.
- For example, since `cyl` and `disp` are highly correlated, and one can find a nearly linear combination between the two.
- Or, it would mean that each observation may be characterized by a single variable, which is a linear combination of `cyl` and `disp`.

Dimension reduction

- Since the original two variables, `cyl` and `disp`, can be used to generate a single variable through their linear combination without necessarily hurting represent-ability, now the originally two variables are reduced to a single variable.
- Thus, PCA is one of the *dimension reduction* techniques, and maybe useful especially if
 - *i)* variables are highly correlated to each other and/or
 - *ii)* there are too many variables to consider.

Preparing covariance matrix

- Remind that, from your elementary statistics course, one variable's sample variance calculated as $\frac{\sum_{i=1}^n (X_i - \bar{X})^2}{n-1}$.
- In order words, from original observations, i) subtract its sample mean (a.k.a. *de-mean* process), then ii) take squared value, and then divide it by $n - 1$.
- Preparing sample covariance matrix is analgous to that as follows.
- Please attempt to understand the following code.

```
# de-meaning process of $X$
# (for each column, x-mean(x) is applied)
X_demeaned <- apply(X, 2, function(x) x-mean(x))
```

- Remind that, for a matrix A , $t(X) \%*\% X$ is equivalent to squaring it.

```
Cov <- (t(X_demeaned) \%*\% X_demeaned)/(nrow(X)-1)
```

- Covariance matrix Cov is a 6×6 matrix, because six variables are involved here.

Cov

```
##          cyl          disp          hp          drat          wt          qsec
## cyl      3.1895161    199.66028    101.93145   -0.66836694    1.3673710   -1.88685484
## disp    199.6602823   15360.79983   6721.15867  -47.06401915   107.6842040  -96.05168145
## hp      101.9314516   6721.15867   4700.86694  -16.45110887    44.1926613  -86.77008065
## drat    -0.6683669   -47.06402   -16.45111    0.28588135   -0.3727207    0.08714073
## wt       1.3673710    107.68420    44.19266   -0.37272073    0.9573790   -0.30548161
## qsec    -1.8868548   -96.05168   -86.77008    0.08714073   -0.3054816    3.19316613
```

Positive semi-definite (psd)

- A proper covariance matrix is positive semi-definite, meaning that its eigenvalues are all non-negative. (Just like your familiar variance is always non-negative)
- Find its eigenvalues and see if they are all non-negative numbers.

```
eigen(Cov)$values
```

```
## [1] 1.861324e+04 1.453791e+03 1.586537e+00 4.584231e-01 1.248216e-01
## [6] 8.738486e-02
```

- All looks good, so we can proceed.

Principal components

- Remind that the six principal components correspond to v_1, v_2, \dots, v_6 .
It can be retrieved as following.

```
eigen(Cov)$vectors
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] -0.012042822 -0.003376004  0.225464493  0.913153920  0.116301042
## [2,] -0.900263018  0.435130058  0.004337984 -0.012130860  0.004694063
## [3,] -0.435075671 -0.899991800 -0.026532118 -0.002512279  0.003469759
## [4,]  0.002661502 -0.003901248  0.057414058 -0.323952312 -0.131060362
## [5,] -0.006242691  0.004868276 -0.203604150  0.191625494 -0.958484375
## [6,]  0.006676402  0.025025499 -0.950627144  0.155984033  0.224879700
##           [,6]
## [1,]  0.318799814
## [2,]  0.000566658
## [3,] -0.001271341
## [4,]  0.935178808
## [5,] -0.055408055
## [6,]  0.143997516
```


Explanation power of each principal component.

- Let the eigenvalues of the matrix Cov be $\lambda_1, \lambda_2, \dots, \lambda_6$ in a descending order.
- The amount of overall variation that the first principal component can explain is $\frac{\lambda_1}{\sum_{i=1}^6 \lambda_i}$. (Check this number is about 92.7%)
- The amount of overall variation that the second principal component can explain is $\frac{\lambda_2}{\sum_{i=1}^6 \lambda_i}$, and so on.
- In sum, the following is the explanation power of each principal components.

```
exp_power <- eigen(Cov)$values/sum(eigen(Cov)$values)
exp_power
```

```
## [1] 9.274490e-01 7.243857e-02 7.905297e-05 2.284202e-05 6.219529e-06
## [6] 4.354157e-06
```

- Cumulatively, the first n principal components' explanation power is:

```
cumsum(exp_power) # cumulative sum
```

```
## [1] 0.9274490 0.9998875 0.9999666 0.9999894 0.9999956 1.0000000
```

"Optimism is the faith that leads to achievement - Hellen Keller"